

Memory & I/O

Yi Kuo @ HPC-I



Outline

I. Memory

- A. Memory Hierarchy
- B. Cache
- C. Pages

II. I/O

- A. Chipset
- B. PCIe
- C. Networking
- D. NVIDIA GPUs
- E. Storage




Memory




Memory Hierarchy

- CPUs could only access registers & load from main memory (DRAM)
- Main Memory is slow (compared to CPU)
 - Every memory access instructions would stall the CPU
 - We need cache!
 - Smaller, Faster, More expensive
- Main Memory is small & expensive (compared to disk)
 - We need to store data on secondary storage!
 - Bigger, Slower, Cheaper

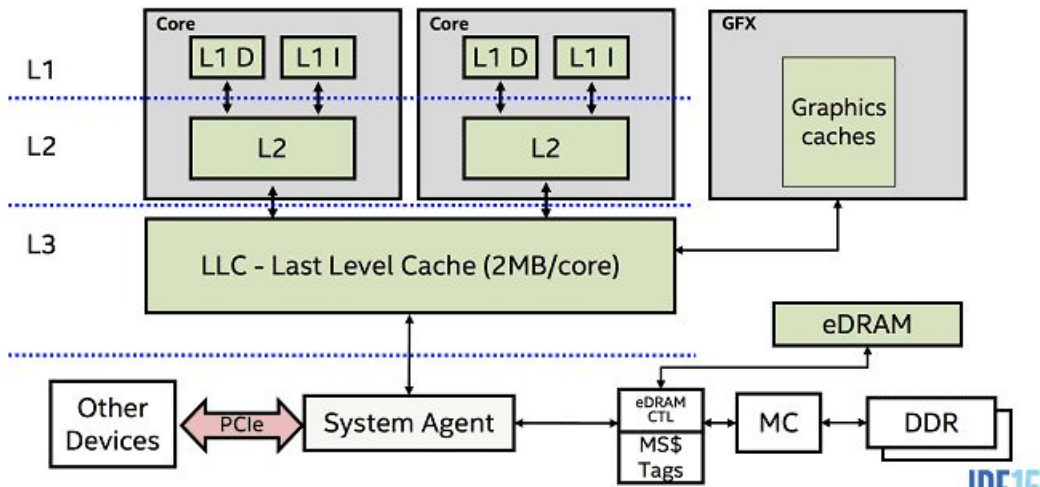


Memory Hierarchy



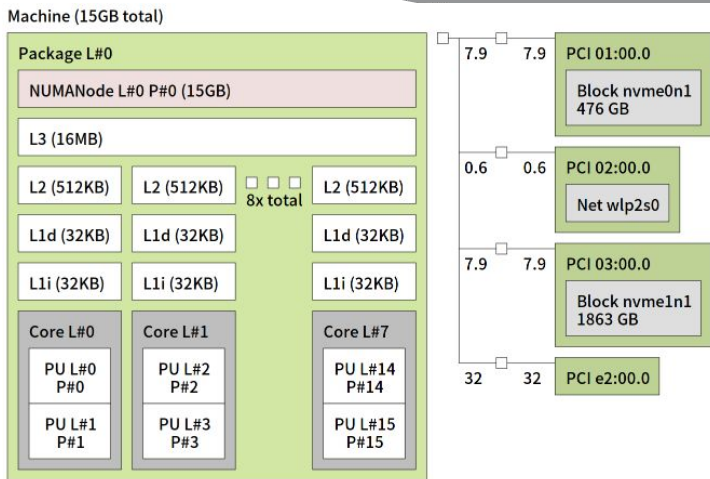
CPU Caches

- Register: 1 cycle
- L1 Cache: 4 cycles
 - L1i: Instruction Cache
 - L1d: Data Cache
- L2 Cache: 12 cycles
- L3 Cache (LLC): ~21 cycles
 - Shared among multiple cores
- RAM: ~117 cycles
- Disk: 10000+ cycles



lstopo

- GUI mode
 - Requires X forwarding (ssh -Y)
- Console mode
 - Without \$DISPLAY variable, or **lstopo-no-graphics**



```
> lstopo-no-graphics
Machine (15GB total)
Package L#0
  NUMANode L#0 (P#0 15GB)
  L3 L#0 (16MB)
    L2 L#0 (512KB) + L1d L#0 (32KB) + L1i L#0 (32KB) + Core L#0
      PU L#0 (P#0)
      PU L#1 (P#1)
    L2 L#1 (512KB) + L1d L#1 (32KB) + L1i L#1 (32KB) + Core L#1
      PU L#2 (P#2)
      PU L#3 (P#3)
    L2 L#2 (512KB) + L1d L#2 (32KB) + L1i L#2 (32KB) + Core L#2
      PU L#4 (P#4)
```

```
HostBridge
PCIBridge
  PCI 01:00.0 (NVMEExp)
    Block(Disk) "nvme0n1"
PCIBridge
  PCI 02:00.0 (Network)
    Net "wlp2s0"
PCIBridge
  PCI 03:00.0 (NVMEExp)
    Block(Disk) "nvme1n1"
PCIBridge
  PCI e2:00.0 (VGA)
```



Cache Line & Locality

- The size of each data RAM entry is referred to “cache line size” or “block size”
- Common cache line size is 64 Bytes
 - Each time CPU loads 64 Bytes from main memory to cache
 - 16x consecutive 32-bit integers are loaded into the cache together
- Spatial Locality
 - Data that is located near each other in memory have a higher probability to be accessed together.
 - Arrays, Structures
- Temporal Locality
 - Data that is accessed repeatedly in a short duration of time is more likely to be accessed again in the near future.
 - Loops, Variables



Lab (as your Homework)

- Matrix Sum

- https://gist.githubusercontent.com/YiPrograms/312b82fc3047ea121c697932b6b289f1/raw/mat_sum.cpp
- `g++ mat_sum.cpp -o mat_sum`

- What's wrong with this program?

- `perf stat -d ./mat_sum`

- Improve it according to cache locality

- Submission:

<https://docs.google.com/document/d/1S2Se8gpPpxFviomCSNQWC0ppGf8CKR6NME3sunRjFE0/edit?usp=sharing>

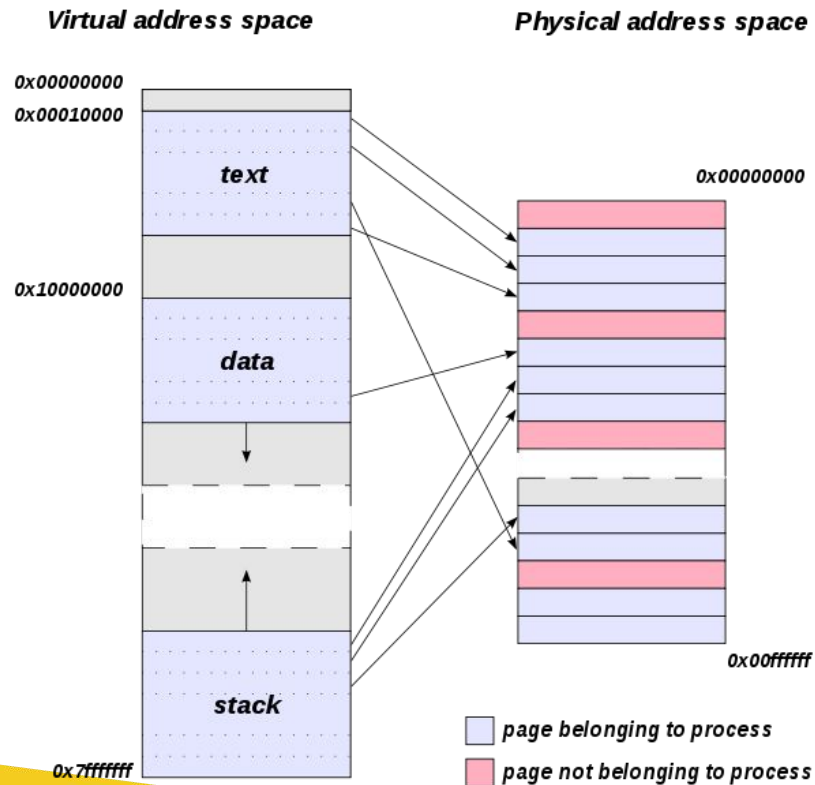


Virtual Memory

- The main memory is a big array
- Each process would want to have contiguous memory
 - -> Fragmentation
- A process should not be able to access the memory of other processes
- Virtual Memory
 - Each process has its own memory address space (Virtual Address, or VA)
 - The instructions in executables use virtual addresses
 - OS stores a **Page Table** to map from virtual address to physical address
 - Every process has a page table
 - The page size is 4KB on most computers
 - **MMU hardware** reads the page table to do the translation
 - Page table in memory -> Slow -> **TLB** Cache




Virtual Memory



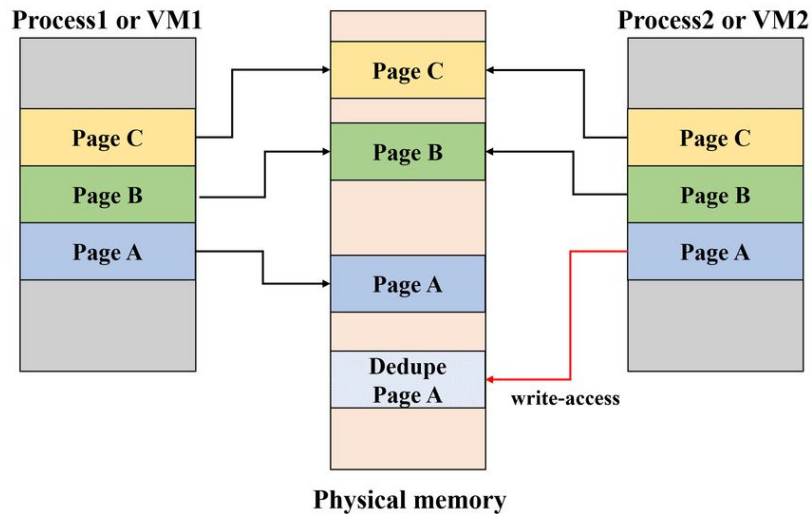
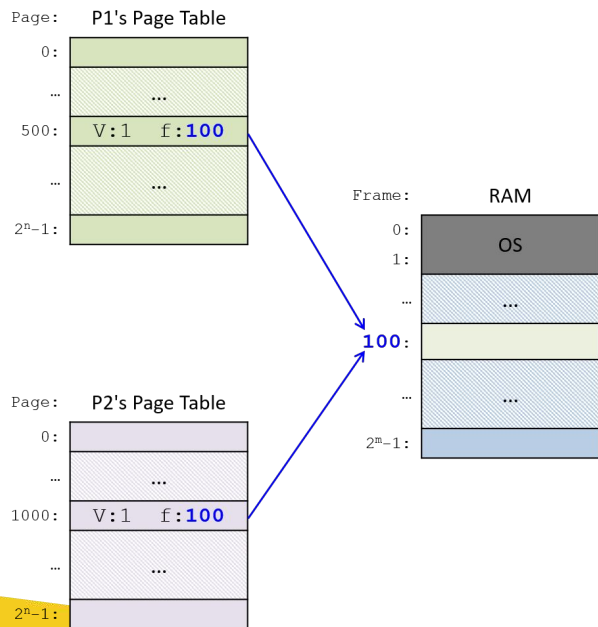
Swapping

- Pages could be moved to secondary storage (e.g. Disks) to save main memory
- When the MMU failed to do the translation (the page is not in the memory, illegal access, etc.), **Page Fault** happens, and the control is handed to the OS
 - If the page is moved to secondary storage
 - OS moves the page back to main memory
 - Continue the program, and MMU would try the translation again
 - If it's an illegal access (e.g. page non-existent, writing read-only pages)
 - OS terminates the program with SIGSEGV (Segmentation Fault)



Shared Pages

- A physical page could be shared among multiple processes (page tables)
 - e.g. Shared Library, Copy on Write

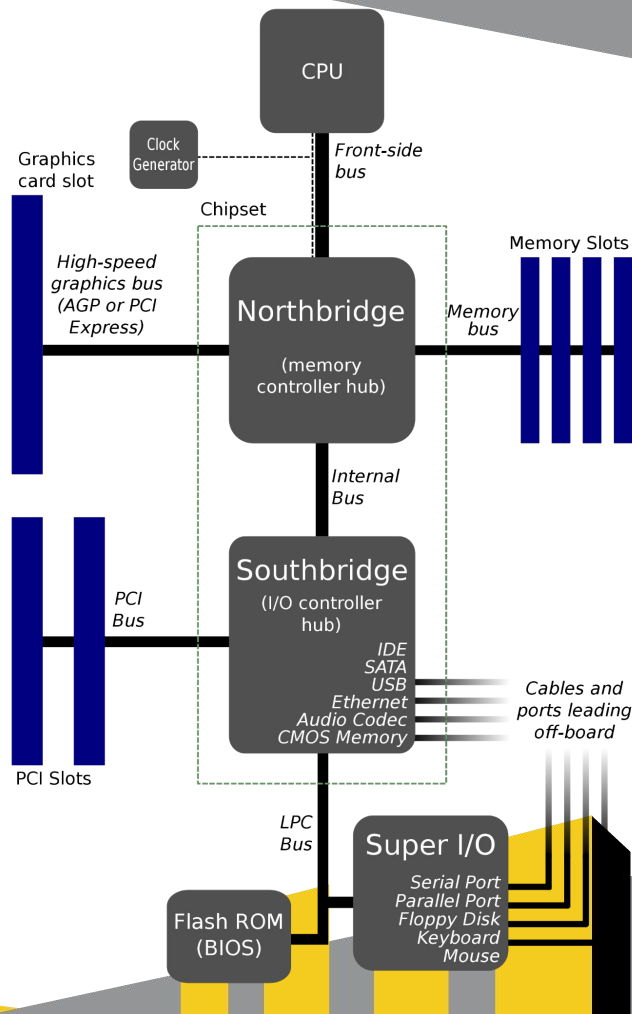


I/O



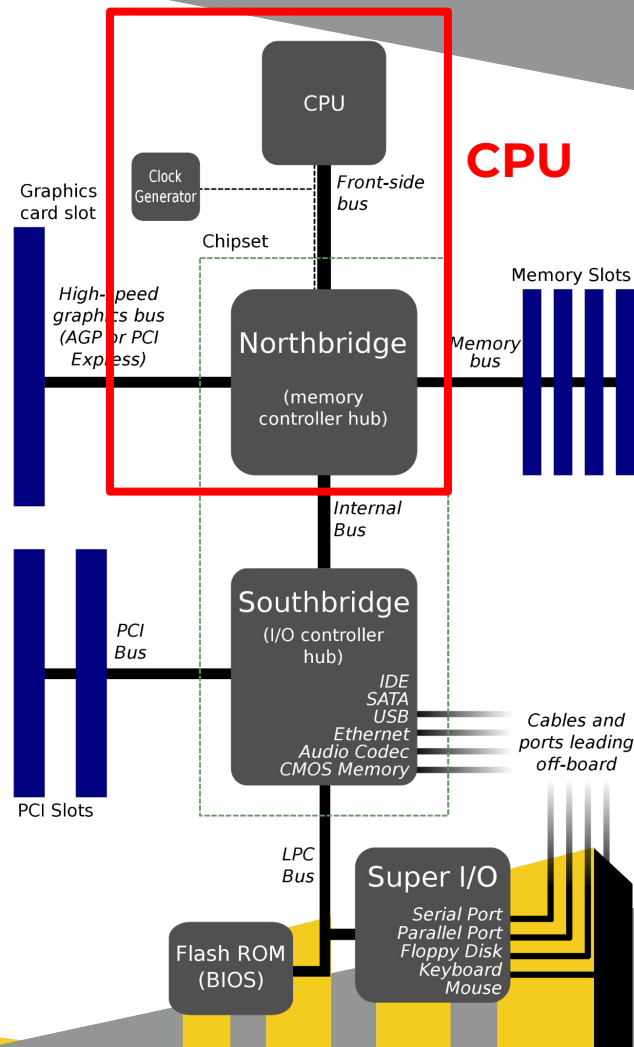
Chipset

- North Bridge
 - Memory controller
 - High speed I/O (e.g. PCIe)
- South Bridge (I/O Controller)
 - Low speed I/O
 - PCI
 - USB
 - SATA
 - etc.



Chipset

- North Bridge
 - Memory controller
 - High speed I/O (e.g. PCIe)
 - **Merged into CPU**
- South Bridge (I/O Controller)
 - Low speed I/O
 - PCI
 - USB
 - SATA
 - etc.

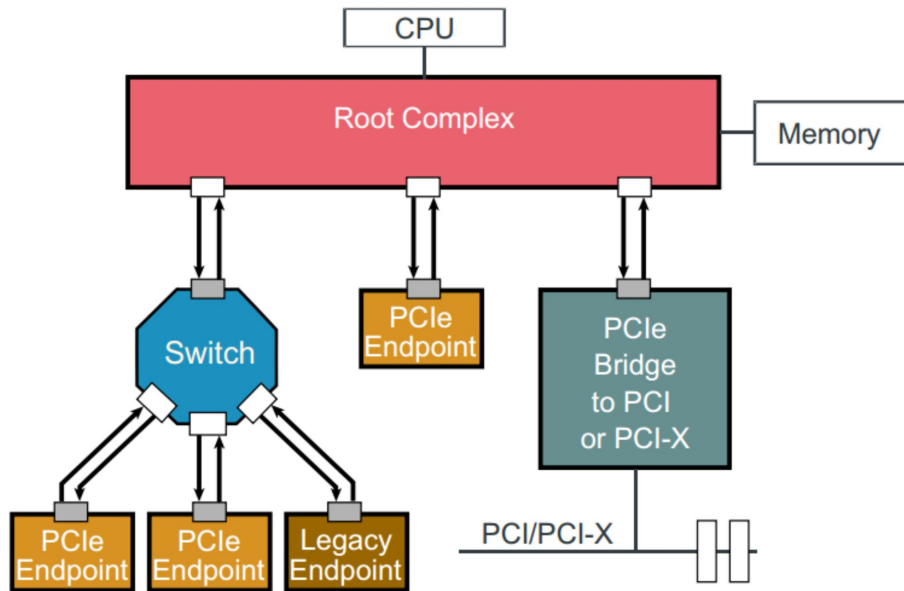


PCIe



PCIe Topology

- Root Complex (RC)
 - The root of PCIe topology
 - Bridges between PCIe & CPU/other components
 - North Bridge
- PCIe Switch
 - Route PCIe packets for multiple PCIe devices
- Endpoint
 - PCIe devices (e.g. GPU, NIC)
 - Native PCIe: Memory Map only
 - Legacy PCIe: support IO requests



PCIe Transactions

- Programmed I/O


- Initiated by the CPU to access PCIe devices

- DMA

- Initiated by the endpoint
- Directly access the memory without CPU's involvement

- P2P


- Initiated from one endpoint, to directly access another endpoint



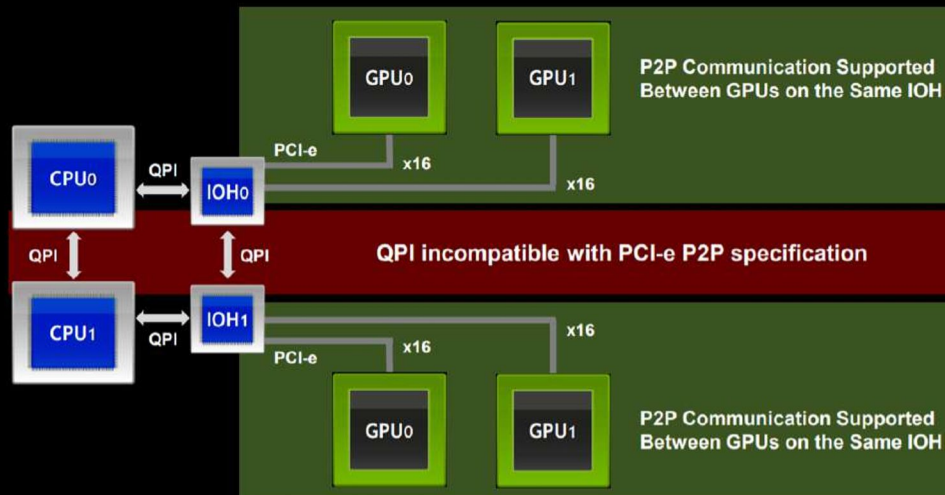
PCIe P2P

- Routing under PCIe Bridges provides the best P2P performance
- Routing via Root Complex would degrade performance (since RC is in the CPU)
- Routing between RCs are not supported by specification

Single IOH Systems with PCI-e Switches: Fully Supported




P2P Communication Across QPI: Not Supported



lstopo

- GUI mode
 - Requires X forwarding (ssh -Y)
- Console mode
 - Without \$DISPLAY variable, or **lstopo-no-graphics**



```
> lstopo-no-graphics
Machine (15GB total)
Package L#0
  NUMANode L#0 (P#0 15GB)
  L3 L#0 (16MB)
    L2 L#0 (512KB) + L1d L#0 (32KB) + L1i L#0 (32KB) + Core L#0
      PU L#0 (P#0)
      PU L#1 (P#1)
    L2 L#1 (512KB) + L1d L#1 (32KB) + L1i L#1 (32KB) + Core L#1
      PU L#2 (P#2)
      PU L#3 (P#3)
    L2 L#2 (512KB) + L1d L#2 (32KB) + L1i L#2 (32KB) + Core L#2
      PU L#4 (P#4)
```

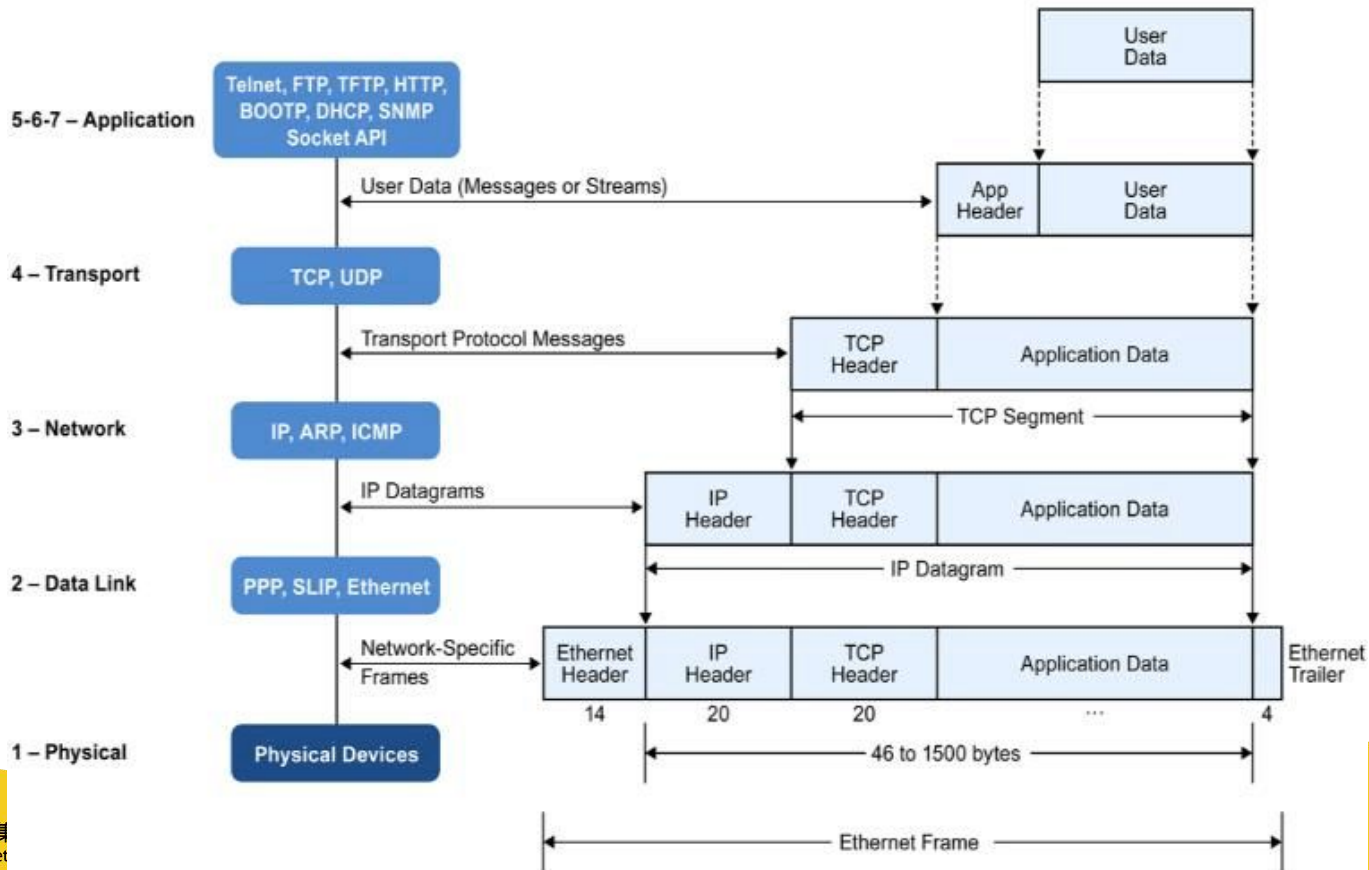
```
HostBridge
PCIBridge
  PCI 01:00.0 (NVMEExp)
    Block(Disk) "nvme0n1"
PCIBridge
  PCI 02:00.0 (Network)
    Net "wlp2s0"
PCIBridge
  PCI 03:00.0 (NVMEExp)
    Block(Disk) "nvme1n1"
PCIBridge
  PCI e2:00.0 (VGA)
```



Networking

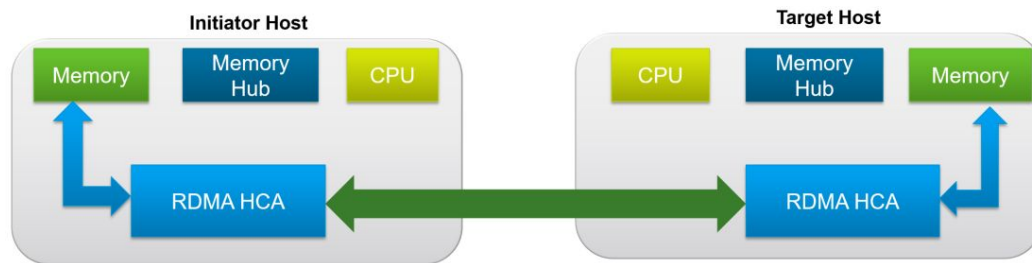


Ethernet (IP Network)

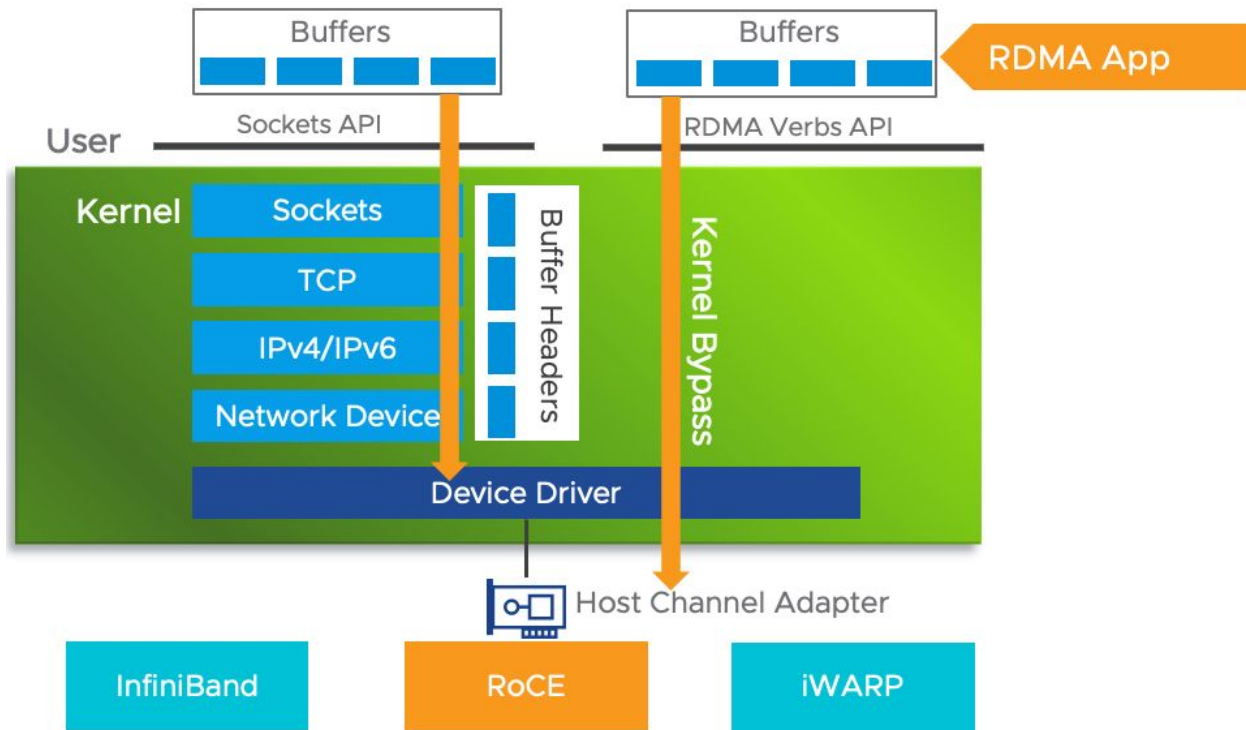


RDMA

- **Remote Direct Memory Access**
- **Hardware Offload**
 - Network protocols for remote memory access and transfer are offloaded to the NIC
- **Kernel Bypass**
 - Without the involvement of Operating System, data can be sent and received directly between buffers of applications without being copied between network layers
- **Protocols**
 - Infiniband
 - RoCE
 - iWARP

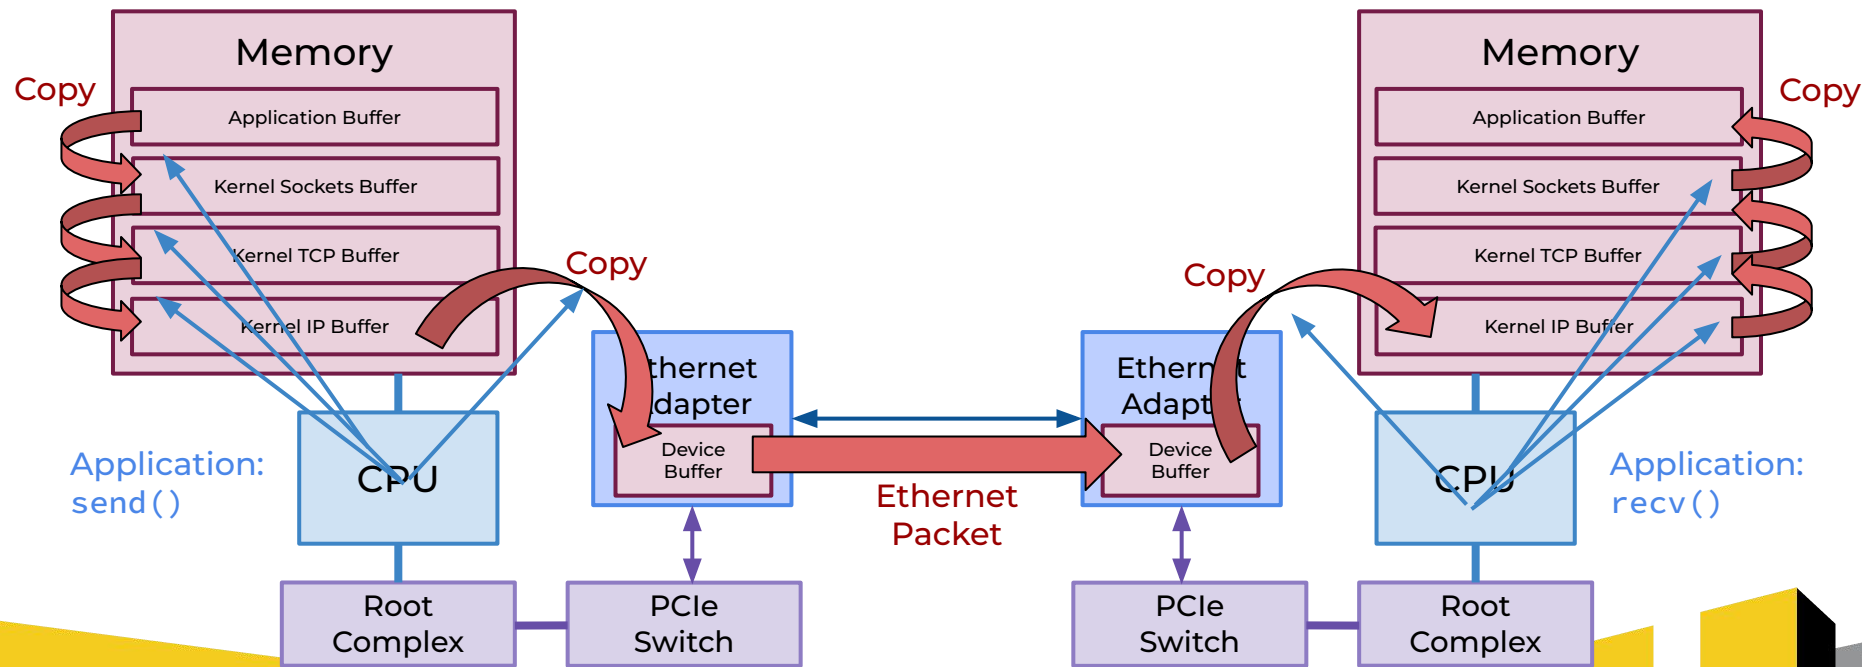


TCP vs RDMA



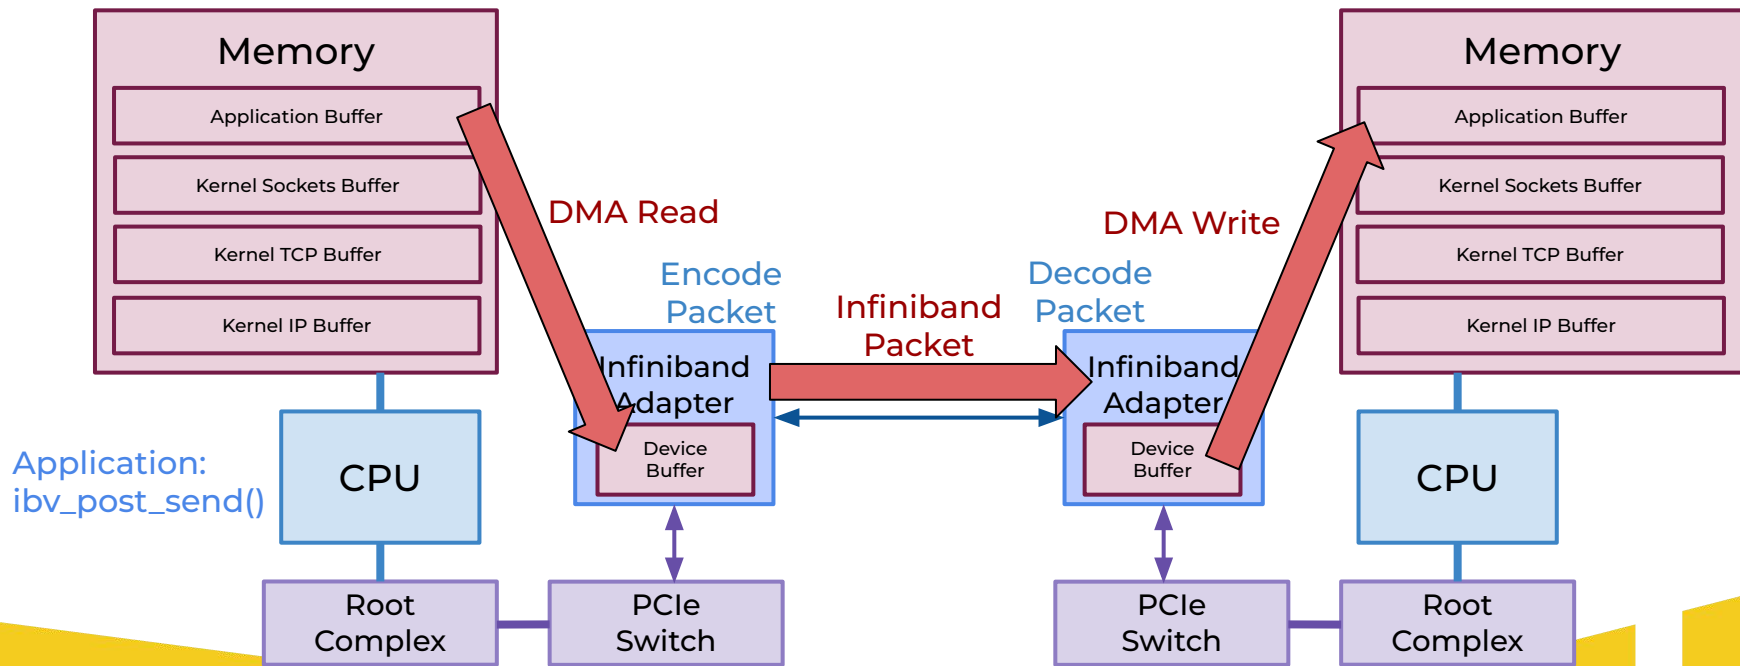
Transferring with Ethernet

- Left sends data to Right

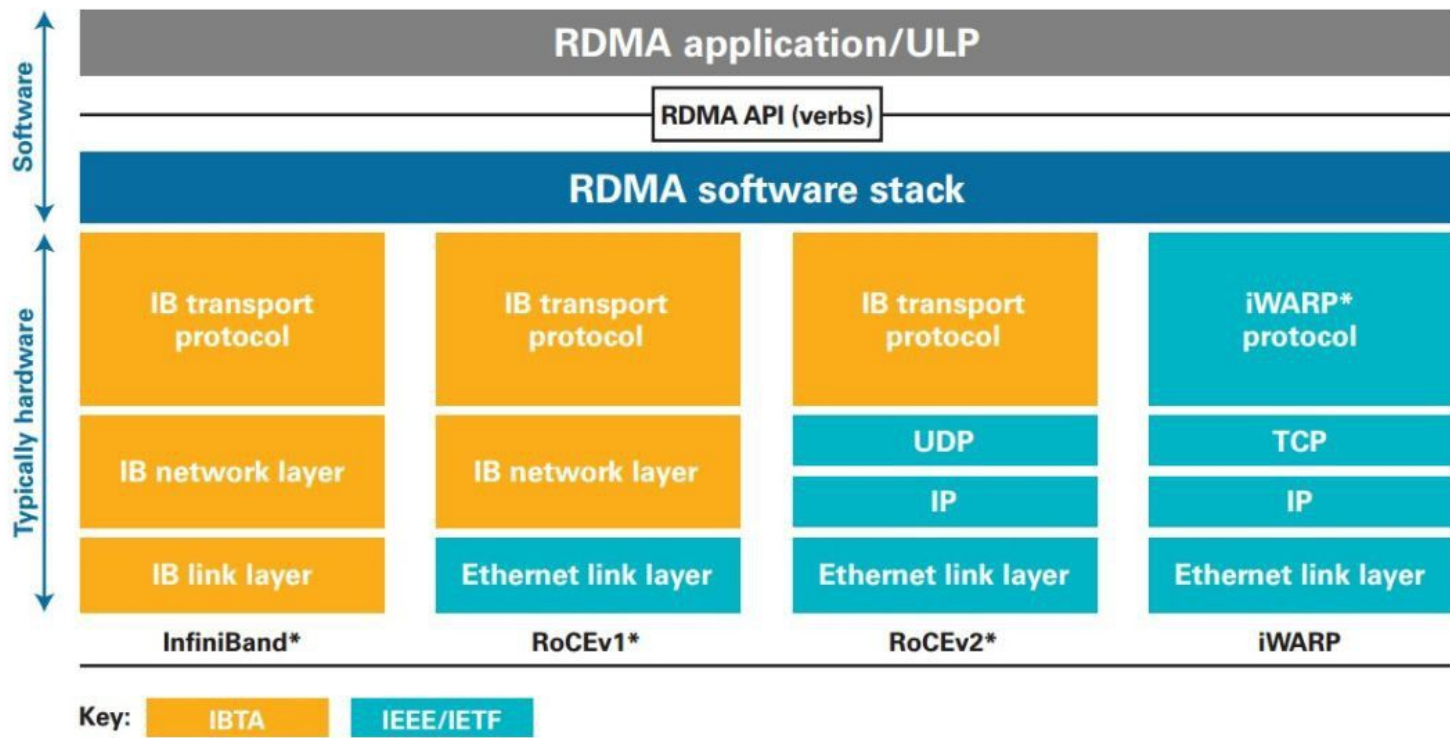


Transferring with RDMA

- Left sends data to Right



RDMA Protocols & Devices



Infiniband Data Rates

Data Rate	Throughput	
	1x	4x (QSFP)
SDR	2 Gb/s	8 Gb/s
DDR	4 Gb/s	16 Gb/s
QDR	8 Gb/s	32 Gb/s
FDR10	10 Gb/s	40 Gb/s
FDR	13.64 Gb/s	54.54 Gb/s
EDR	25 Gb/s	100 Gb/s
HDR	50 Gb/s	200 Gb/s
NDR	100 Gb/s	400 Gb/s
XDR	200 Gb/s	800 Gb/s
GDR	400 Gb/s	1600 Gb/s



QSFP vs SFP



NVIDIA GPUs



NVLINK

- NVLINK provides high-speed P2P communication between GPUs
- PCIe Version's NVLINK could only connect pairs of nearby GPUs
- SXM Version's NVLINK is built on the baseboard
- `nvidia-smi nvlink -R -s`

PCIe Version



SXM Version



nvidia-smi topo -m

```
[yikuo0425@gn1103 ~]$ nvidia-smi topo -m
```

	<u>GPU0</u>	<u>GPU1</u>	<u>GPU2</u>	<u>GPU3</u>	<u>GPU4</u>	<u>GPU5</u>	<u>GPU6</u>	<u>GPU7</u>	<u>mlx5_0</u>	<u>mlx5_1</u>	<u>mlx5_4</u>	<u>mlx5_5</u>	<u>CPU Affinity</u>	<u>NUMA Affinity</u>
GPU0	X	NV1	NV1	NV2	SYS	SYS	NV2	SYS	PIX	NODE	SYS	SYS	0-16	0
GPU1	NV1	X	NV2	NV1	SYS	SYS	SYS	NV2	PIX	NODE	SYS	SYS	0-16	0
GPU2	NV1	NV2	X	NV2	NV1	SYS	SYS	SYS	NODE	PIX	SYS	SYS	0-16	0
GPU3	NV2	NV1	NV2	X	SYS	NV1	SYS	SYS	NODE	PIX	SYS	SYS	0-16	0
GPU4	SYS	SYS	NV1	SYS	X	NV2	NV1	NV2	SYS	SYS	PIX	NODE	18-32	1
GPU5	SYS	SYS	SYS	NV1	NV2	X	NV2	NV1	SYS	SYS	PIX	NODE	18-32	1
GPU6	NV2	SYS	SYS	SYS	NV1	NV2	X	NV1	SYS	SYS	NODE	PIX	18-32	1
GPU7	SYS	NV2	SYS	SYS	NV2	NV1	NV1	X	SYS	SYS	NODE	PIX	18-32	1
mlx5_0	PIX	PIX	NODE	NODE	SYS	SYS	SYS	SYS	X	NODE	SYS	SYS		
mlx5_1	NODE	NODE	PIX	PIX	SYS	SYS	SYS	SYS	NODE	X	SYS	SYS		
mlx5_4	SYS	SYS	SYS	SYS	PIX	PIX	NODE	NODE	SYS	SYS	X	NODE		
mlx5_5	SYS	SYS	SYS	SYS	NODE	NODE	PIX	PIX	SYS	SYS	NODE	X		

Legend:

X = Self

SYS = Connection traversing PCIe as well as the SMP interconnect between NUMA nodes (e.g., QPI/UPI)

NODE = Connection traversing PCIe as well as the interconnect between PCIe Host Bridges within a NUMA node

PHB = Connection traversing PCIe as well as a PCIe Host Bridge (typically the CPU)

PXB = Connection traversing multiple PCIe bridges (without traversing the PCIe Host Bridge)

PIX = Connection traversing at most a single PCIe bridge

NV# = Connection traversing a bonded set of # NVLinks



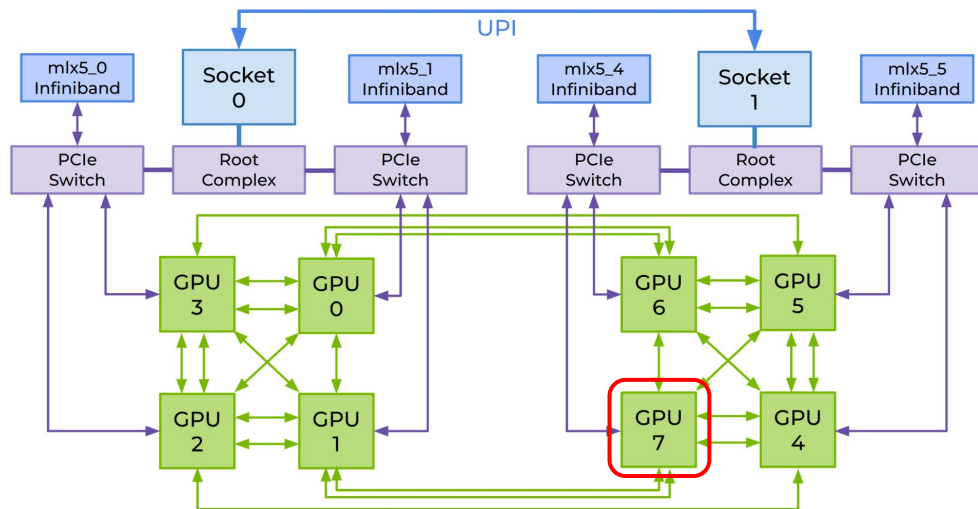
nvidia-smi topo -m

```
[yikuo0425@gn1103 ~]$ nvidia-smi topo -m
```

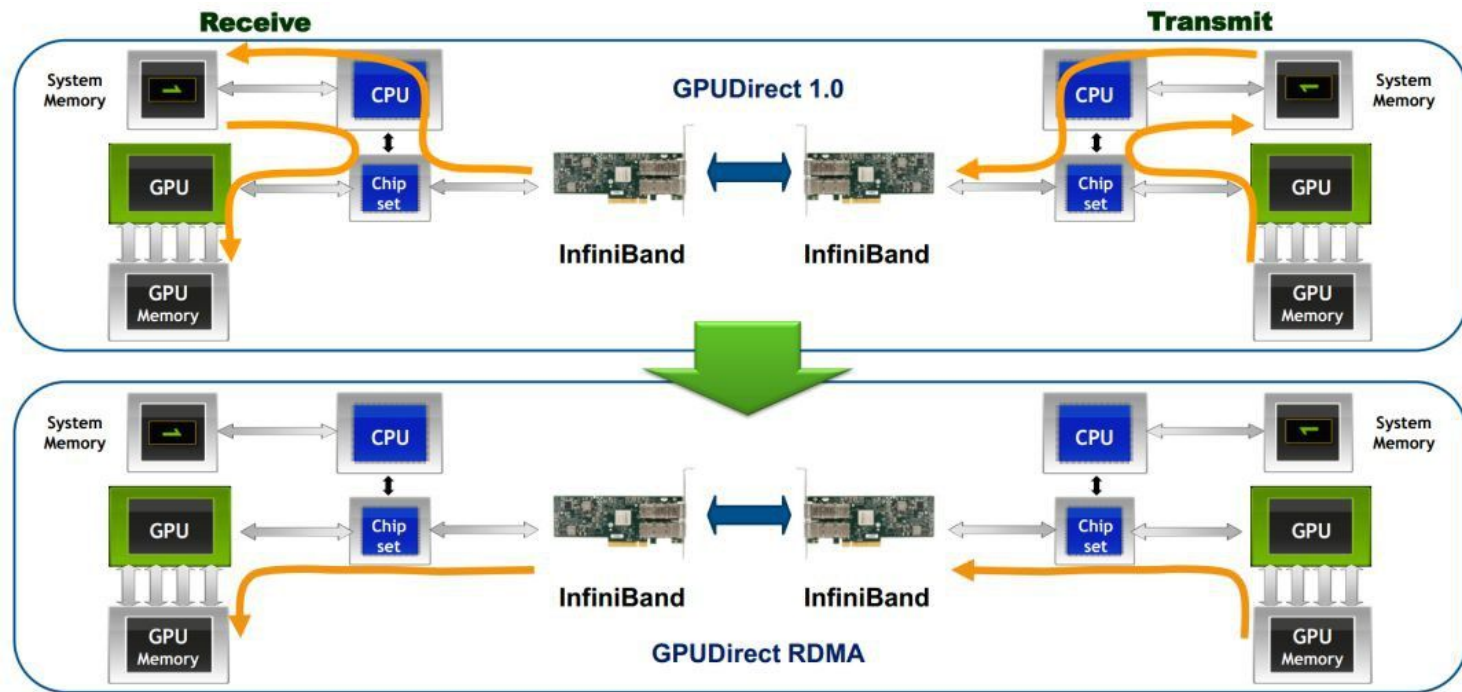
	GPU0	GPU1	GPU2	GPU3	GPU4	GPU5	GPU6	GPU7	mlx5_0	mlx5_1	mlx5_4	mlx5_5	CPU Affinity	NUMA Affinity
GPU0	X	NV1	NV1	NV0	SYS	SYS	NV2	SYS	PIX	NODE	SYS	SYS	0-16	0
GPU6	NV2	SYS	SYS	SYS	NV1	NV2	X	NV1	SYS	SYS	NODE	PIX	18-32	1
GPU7	SYS	NV2	SYS	SYS	NV2	NV1	NV1	X	SYS	SYS	NODE	PIX	18-32	1
mlx5_0	PIX	PIX	NODE	NODE	SYS	SYS	SYS	SYS	X	NODE	SYS	SYS		
mlx5_1	NODE	NODE	SYS	SYS	SYS	SYS	SYS	SYS	NODE	X	SYS	SYS		
mlx5_4	SYS	SYS	SYS	SYS	SYS	SYS	SYS	SYS	SYS	SYS	X	SYS		
mlx5_5	SYS	SYS	SYS	SYS	SYS	SYS	SYS	SYS	SYS	SYS	SYS	X		

GPU7 is connected to the PCIe RC on NUMA node 1, and it would communicate to

- GPU 0, 2, 3
 - UPI over NUMA nodes
- GPU 5, 6
 - A pair of NVLINK
- GPU 1, 4
 - Two pairs of NVLINK
- mlx5_0 and mlx5_1 (IB NIC)
 - UPI over NUMA nodes
- mlx5_4
 - PCIe RC on the same NUMA node
- mlx5_5
 - PCIe Switch



GPUDirect RDMA



- Ensure that the HCA is on the same PCIe Switch as the GPU for best performance

Storage



HDD & SSD


- HDD (Hard Disk Drive)

- Mechanical Disks
- Cheap
- Slow

- SSD (Solid State Drive)


- Flash Chips
- Not so cheap
- Fast

HDD
3.5"



Shock resistant up to 55g (operating)
Shock resistant up to 350g (non-operating)

SSD
2.5"



Shock resistant up to 1500g
(operating and non-operating)

Protocols & Interfaces

- IDE (PATA)





- SATA

- SAS

- SAS drives are mostly used on servers, and has higher RPM.
- SATA drives can be installed in SAS socket, but not vice versa.


- NVMe (PCIe)

- PCIe
- M.2 / NGFF (M Key)
- U.2 / U.3
 - SATA & SAS drives can be installed in U.2 socket, but not vice versa.



U.2 (formerly SFF-8639) (PCIe/NVMe)

■ PCIe DATA, REFERENCE CLOCK AND SIDE BAND



RAID


- Redundant Array of Independent Disks
 - Combining multiple disks to create a big storage, for speed & redundancy

BREAKDOWN OF COMMON RAID LEVELS


Hewlett Packard Enterprise

RAID LEVEL	METHOD	HARDWARE / SOFTWARE	MINIMUM # OF DISKS	COMMON USAGE	PROS	CONS
JBOD	SPANNING		2	INCREASE CAPACITY	COST-EFFECTIVE STORAGE	NO PERFORMANCE OR SECURITY BENEFITS
0	STRIPING		2	HEAVY READ OPERATIONS	HIGH PERFORMANCE (SPEED)	DATA IS LOST IF ONE DISK FAILS
1	MIRRORING		2	STANDARD APP SERVERS	FAULT TOLERANCE, HIGH READ PERFORMANCE	LAG FOR WRITE OPS, REDUCED STORAGE (BY 1/2)
5	STRIPING & PARITY		3	NORMAL FILE STORAGE & APP SERVERS	SPEED + FAULT TOLERANCE	LAG FOR WRITE OPS, REDUCED STORAGE (BY 1/3)
6	STRIPING & DOUBLE PARITY		4	LARGE FILE STORAGE & APP SERVERS	EXTRA LEVEL OF REDUNDANCY, HIGH READ PERFORMANCE	LOW WRITE PERFORMANCE, REDUCED STORAGE (BY 2/5)
10 (1+0)	STRIPING & MIRRORING		4	HIGHLY UTILIZED DATABASE SERVERS	WRITE PERFORMANCE + STRONG FAULT TOLERANCE	REDUCED STORAGE (1/2), LIMITED SCALABILITY


RAID 0



RAID 1

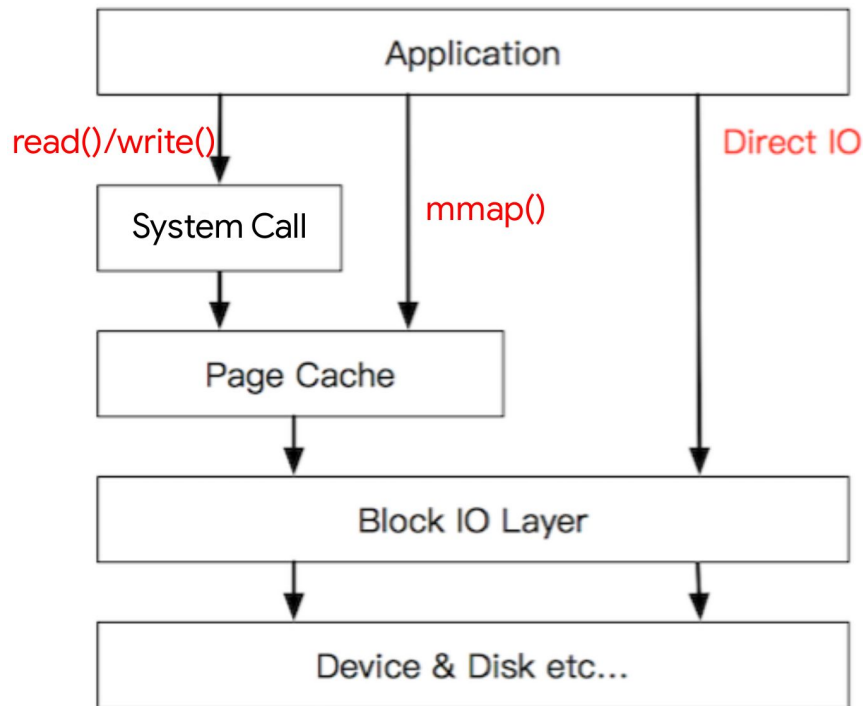


RAID 5



Linux File I/O

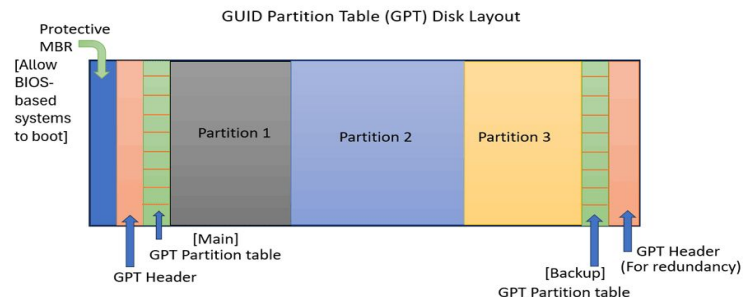
- Buffered I/O
 - Files are buffered with page caches
 - read() / write()
 - fread() / fwrite()
 - With buffer in user space
 - mmap()
 - Directly map page caches to user pages
 - Reduce the overhead of system calls
- Direct I/O (O_DIRECT)
 - Skip page caches
 - Read/Write directly from/to disk
 - Used in Database systems, which manage caches itself



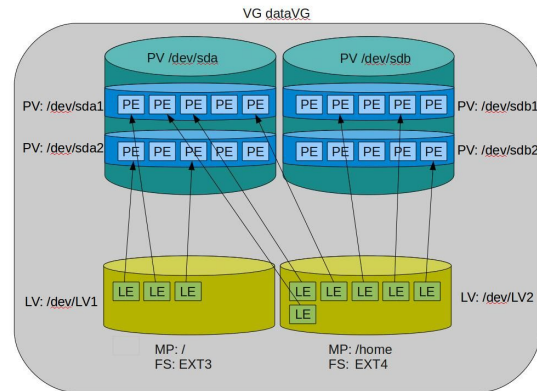
Partitions & LVM

- A disk is partitioned into multiple partitions
 - The partition table describes the partitions on the disk

- MBR (Master Boot Record)
 - For legacy BIOS Systems
 - Max 2TB Disk
- GPT (GUID Partition Table)
 - For UEFI Systems



- LVM (Logical Volume Management)
 - Reducing fragmentation (Just like virtual memory)
 - Features
 - Dynamic volume resizing
 - Block level snapshot
 - RAID



File Systems

- Used to store the file & directory structures
- FAT (File Allocation Table)
- Journaling file system
 - e.g. ext3 / ext4, NTFS, HFS+
- ZFS & Btrfs
 - Support many features
e.g. Snapshot, Compression, Encryption
- NFS (Network File System)
- tmpfs (RAM Disk)
 - `sudo mount -t tmpfs tmpfs /mnt`
- Distributed File System
 - e.g. BeeGFS, Lustre, GlusterFS

