

# 平行計算與程式 補充

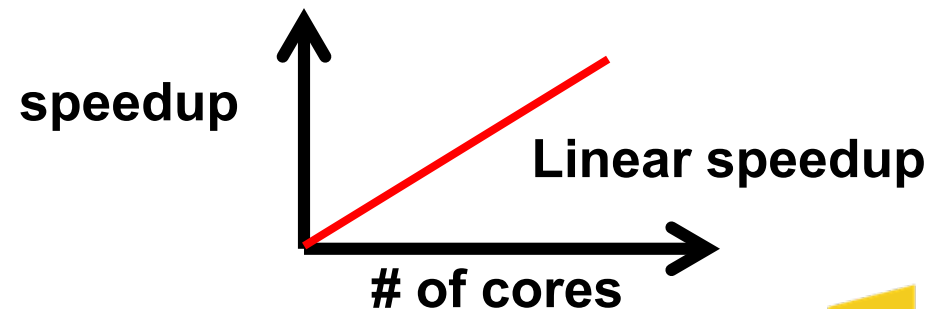
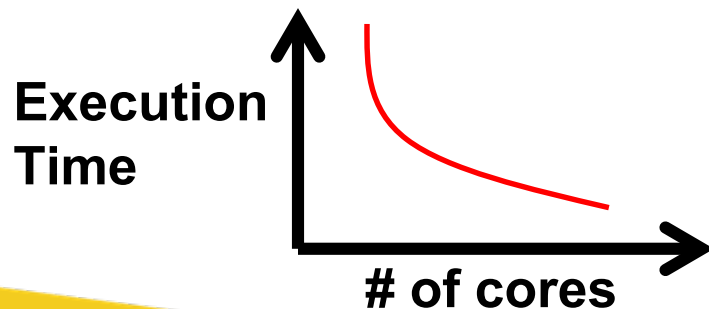


# Strong Scalability vs Weak Scalability



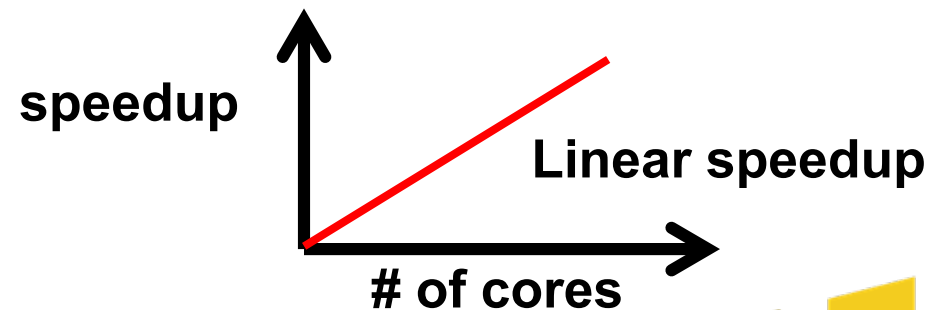
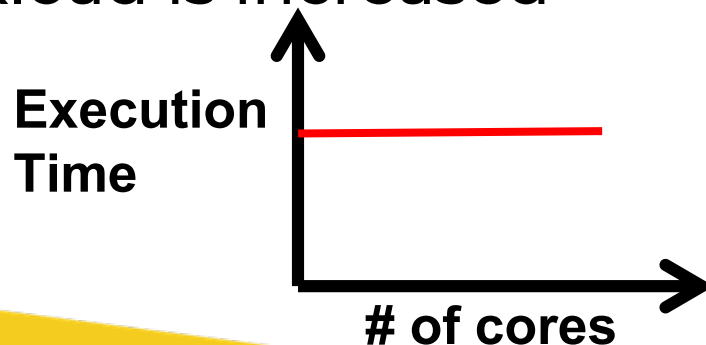
# Strong Scaling

- The **problem size stays fixed** but the number of processing elements are increased.
- It is used to **find a "sweet spot"** that allows the computation to complete in a reasonable amount of time, yet does not waste too many cycles due to parallel overhead.
- Linear scaling is achieved if the speedup is equal to the number of processing elements.



# Weak Scaling

- The **problem size (workload)** assigned to each processing element **stays fixed** and additional processing elements are used to **solve a larger total problem**
- It is a **justification for programs that take a lot of memory** or other system resources (e.g., a problem wouldn't fit in RAM on a single node)
- **Linear scaling** is achieved if the run time stays constant while the workload is increased



# Strong Scaling vs. Weak Scaling

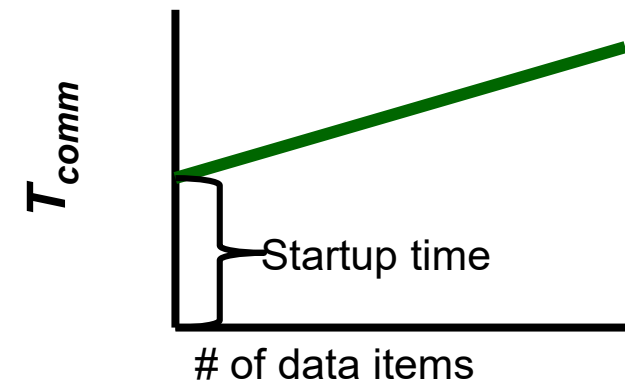
- Strong scaling
  - Linear scaling is harder to achieve, because of the **communication overhead may increase proportional to the scale**
- Weak scaling
  - Linear scaling is easier to achieve because **programs typically employ nearest-neighbor communication patterns** where the **communication overhead is relatively constant** regardless of the number of processes used

# Time Complexity Analysis



# Time Complexity Analysis

- $T_p = T_{comp} + T_{comm}$ 
  - $T_p$ : Total execution time of a parallel algorithm
  - $T_{comp}$ : Computation part
  - $T_{comm}$ : Communication part
- $T_{comm} = q (T_{startup} + n T_{data})$ 
  - $T_{startup}$ : Message latency (assumed constant)
  - $T_{data}$ : Transmission time to send one data item
  - $n$ : Number of data items in a message
  - $q$ : Number of message



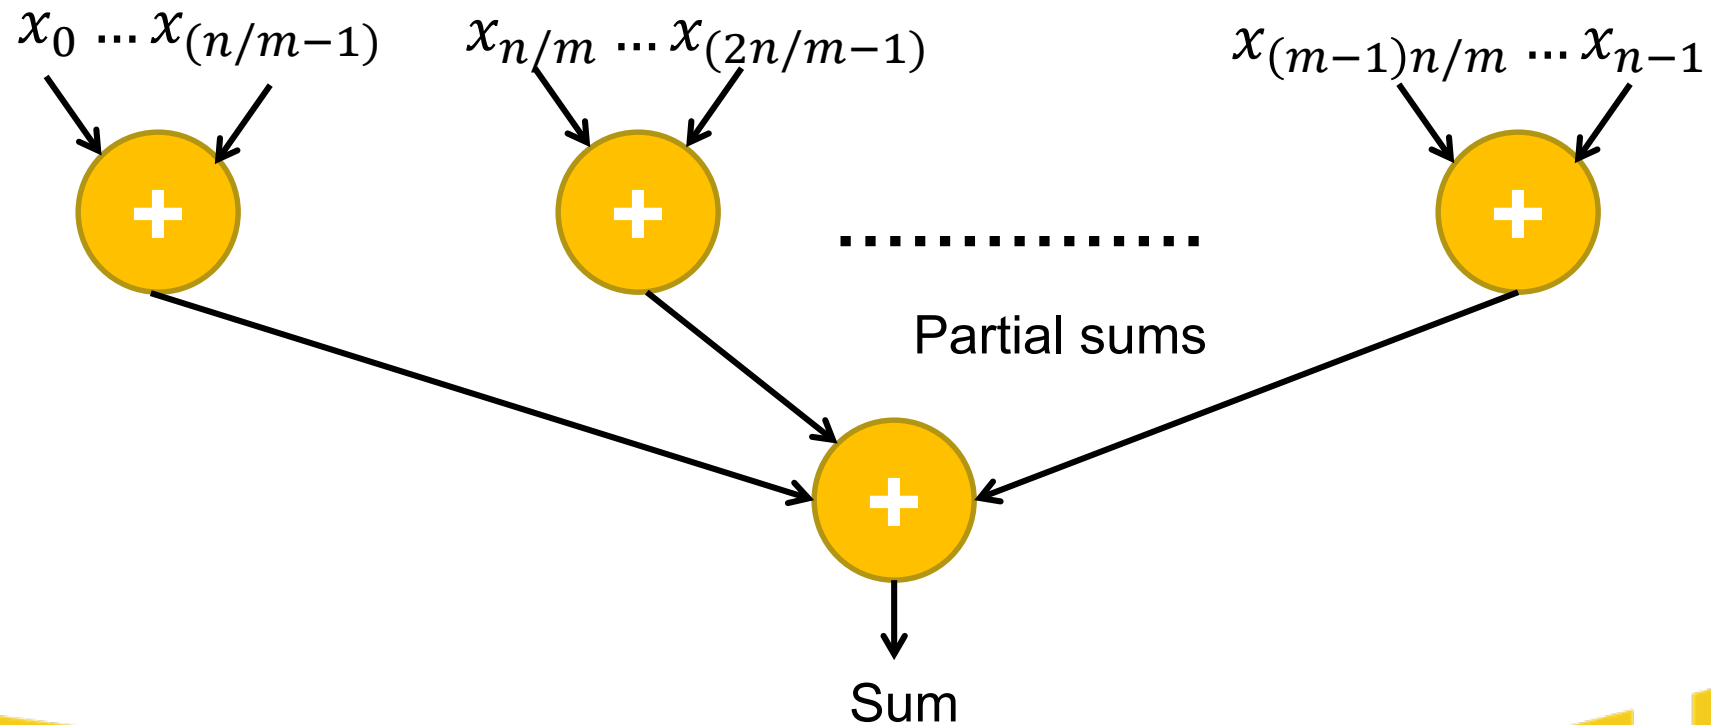
# Time Complexity Example 1

- Algorithm phase:
  1. Computer 1 **sends**  $n/2$  numbers to computer 2
  2. Both computers **add**  $n/2$  numbers simultaneously
  3. Computer 2 **sends** its partial result back to computer 1
  4. Computer 1 **adds** the partial sums to produce the final result
- Complexity analysis:
  - Computation (for step 2 & 4):
    - $T_{\text{comp}} = n/2 + 1 = O(n)$
  - Communication (for step 1 & 3):
    - $T_{\text{comm}} = (T_{\text{startup}} + n/2 \times T_{\text{data}}) + (T_{\text{startup}} + T_{\text{data}})$   
 $= 2T_{\text{startup}} + (n/2 + 1) T_{\text{data}} = O(n)$
  - **Overall complexity:  $O(n)$**



## Time Complexity Example 2

- Adding  $n$  numbers using  $m$  processes
  - Evenly partition numbers to processes



# Time Complexity Example 2

Adding  $n$  numbers using  $m$  processes

- Sequential:  $O(n)$
- Parallel:

- Phase1: Send numbers to slaves

$$t_{comm1} = m(t_{startup} + (n/m)t_{data})$$

- Phase2: Compute partial sum

$$t_{comp1} = n/m - 1$$

- Phase3: Send results to master

$$t_{comm2} = m(t_{startup} + t_{data})$$

- Phase4: Compute final accumulation

$$t_{comp2} = m - 1$$

- Overall:

$$t_p = 2mt_{startup} + (n + m)t_{data} + m + \frac{n}{m} - 2 = O(m + n/m)$$

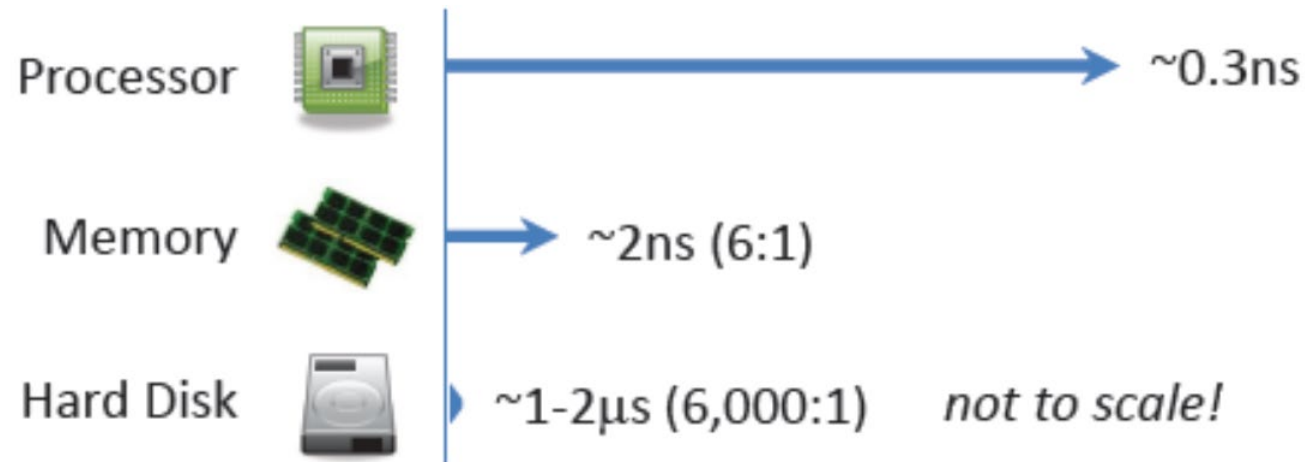
Tradeoff between  
**computation &  
communication**

# MPI & Parallel IO



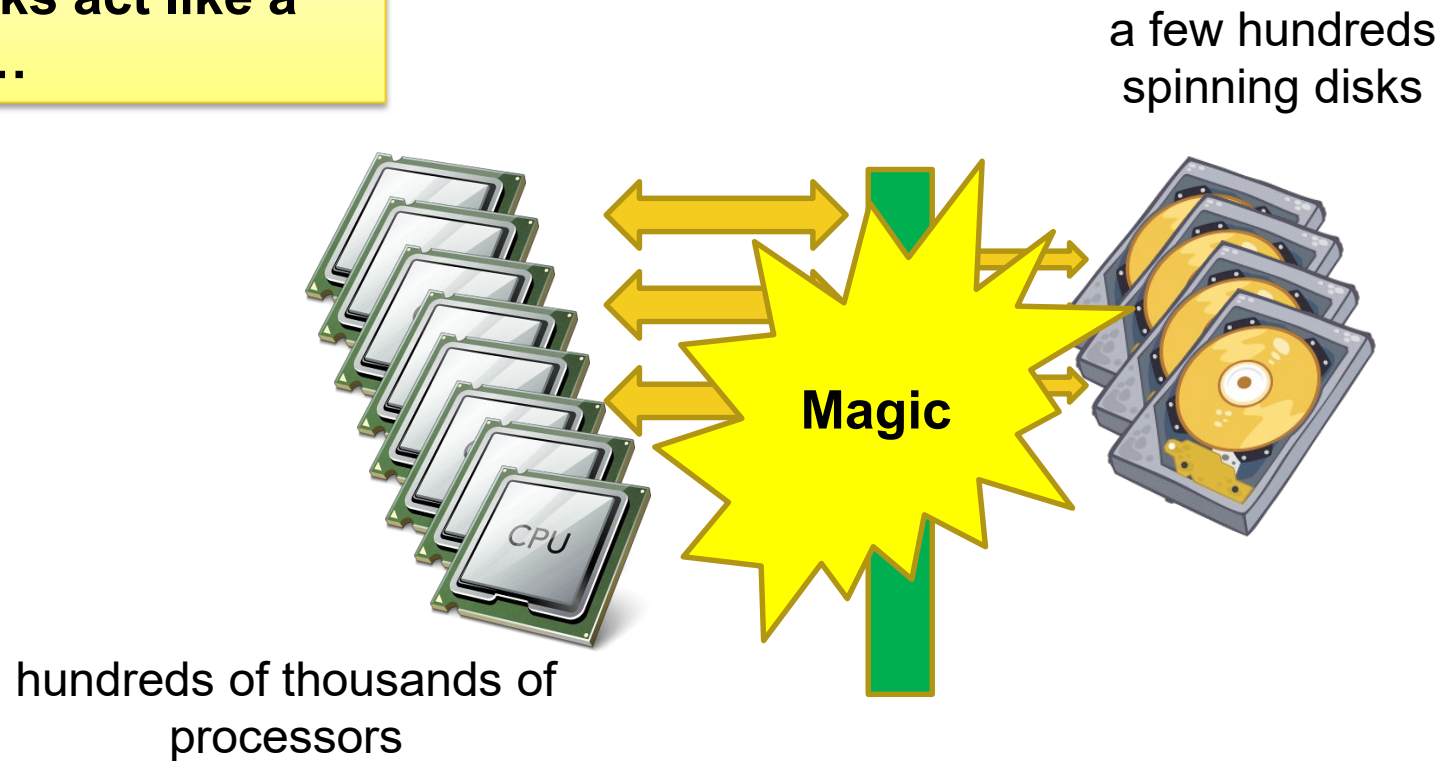
# Relative Speed of Components in HPC Platform

- An HPC platform's I/O subsystems are typically slow as compared to its other parts
- The I/O gap between memory speed and average disk access stands at roughly  $10^{-3}$



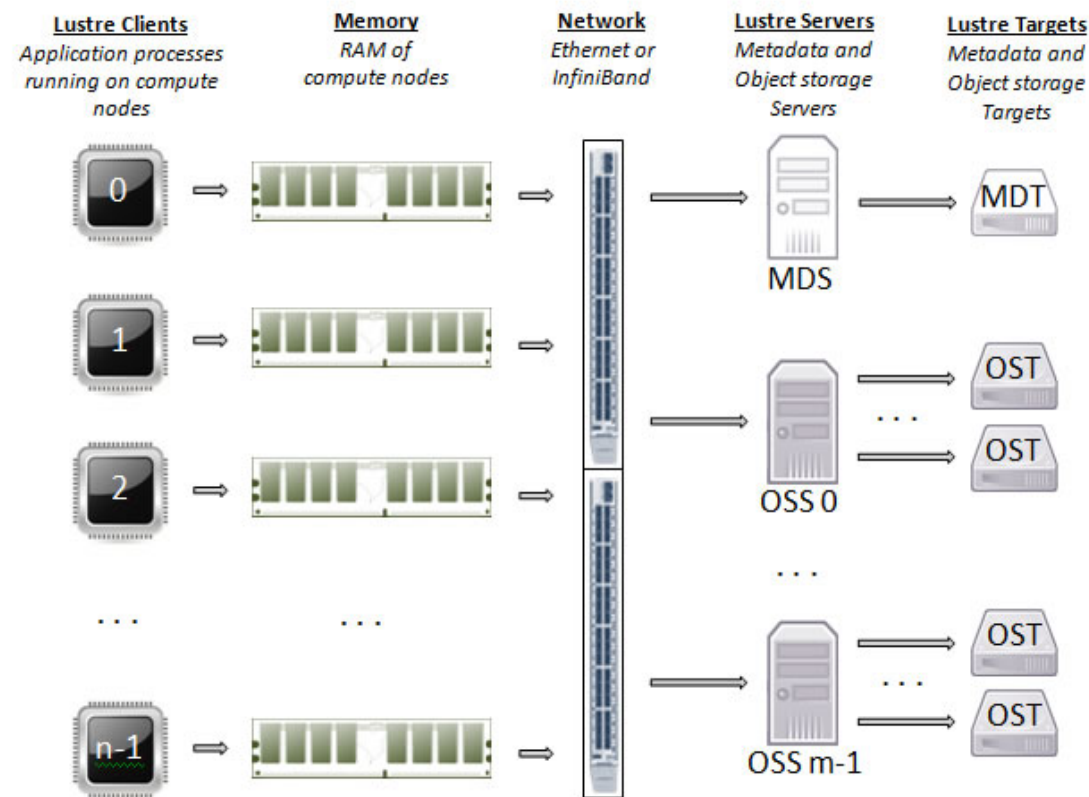
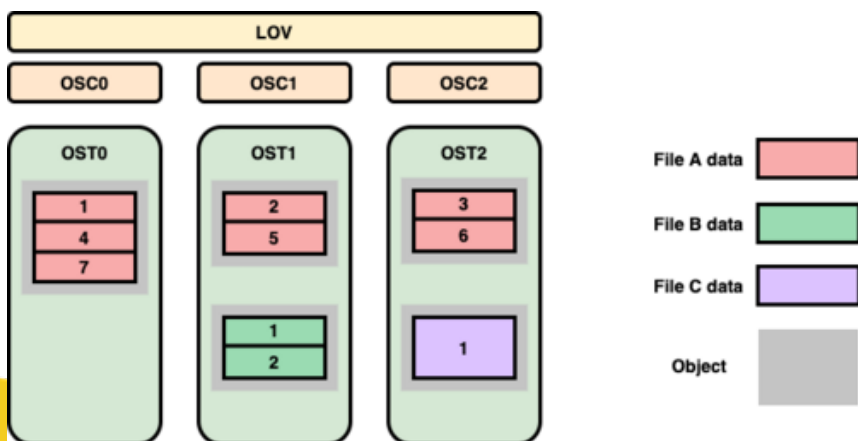
# Concurrent Data Access in a Cluster

**We need some magic to make the collection of spinning disks act like a single disk ...**



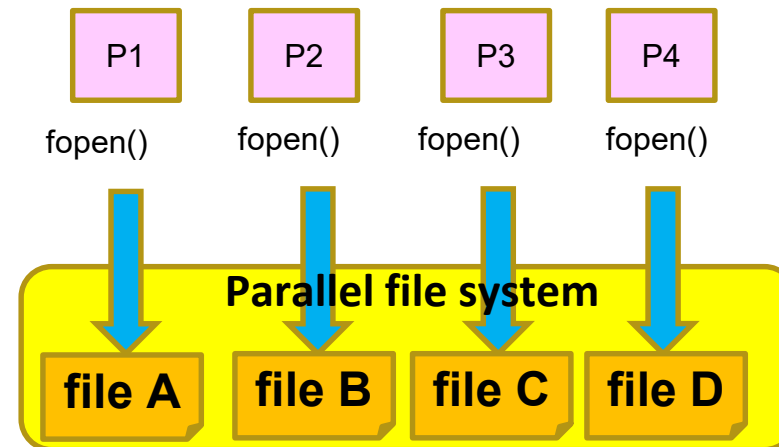
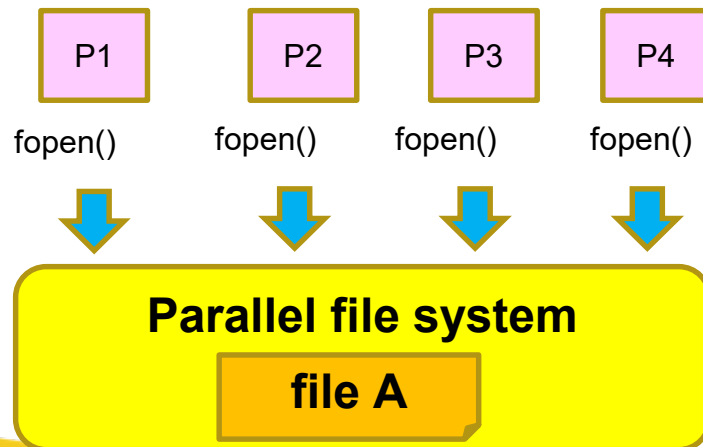
# Parallel File Systems: Lustre

- Separate control plan (metadata) and data plan (data)
- Distributed system architecture
  - Multiple MDS and OSS servers
- Simplify the task of storage node
- Parallel I/O on a single file
  - Files are chunked & striped
  - **Stipe size & count** is configured by users



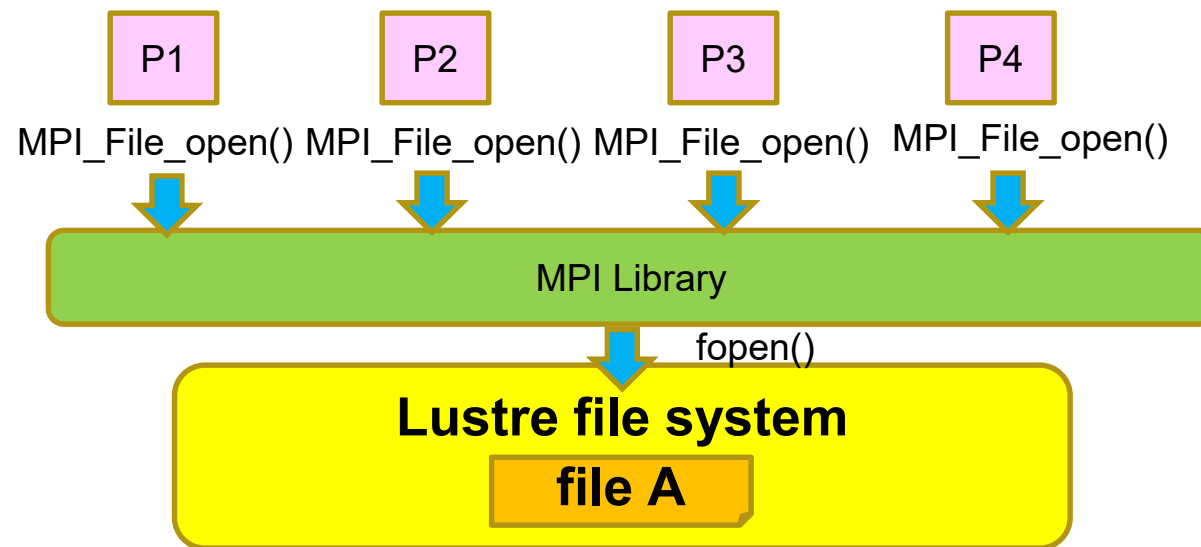
# POSIX File Access Operations

- POSIX file system call “fopen()”:
  - The same is opened by each processes → multiple file handlers across your MPI processes
  - Open the same file with read permission is OK
  - But can't open with **write permission** together due file system **locking** mechanism → data inconsistency
  - To write simultaneously must **create multiple files (can't take advantage of parallel file system & hard to manage)**



# MPI-IO File Access Operations

- MPI-IO call “MPI\_File\_open()”
  - File is **opened only once** in a **collective** manner
  - MPI library will share and synchronize with each other to use the same file handler
  - Can handle both read and write together

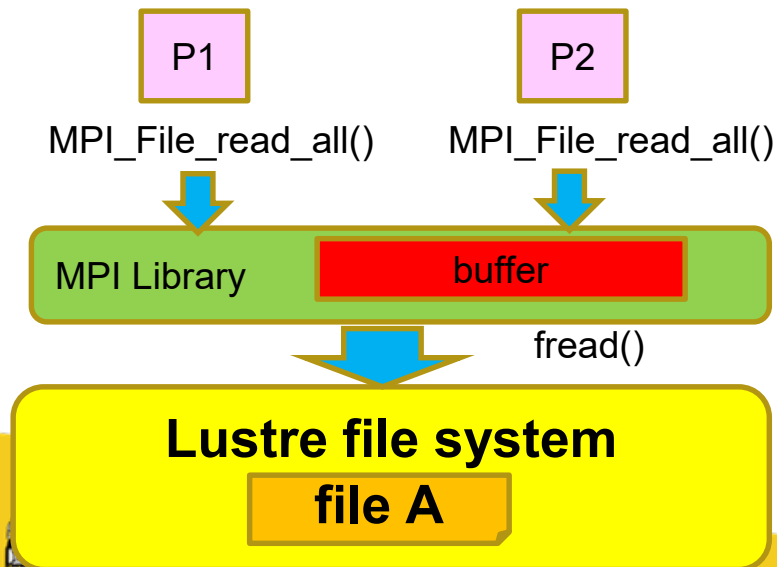




# MPI-IO Independent/Collective I/O

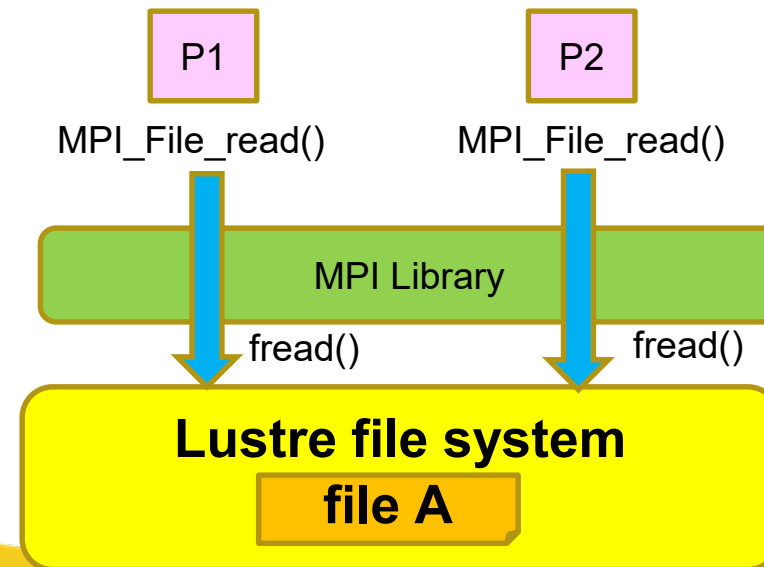
- Collective I/O

- Read/write to a shared memory buffer, then issue **ONE file request**
- Reduce #I/O request  
→ Good for small I/O
- Require **synchronization**



- Independent I/O

- Read/write individually
- Prevent synchronization
- One request per process
- Request is **serialized** if access the same OST
- **Good for large I/O**



# MPI-IO API

- `MPI_File_open(MPI_Comm comm, char *filename, int amode, MPI_Info info, MPI_File *fh)`
  - Open a file
- `MPI_File_close(MPI_File *fh)`
  - Close a file
- `MPI_File_read/write(MPI_File fh, void *buf, int count, MPI_Datatype datatype, MPI_Status *status)`
  - **Independent** read/write using individual file pointer
- `MPI_File_read/write_all(MPI_File fh, void *buf, int count, MPI_Datatype datatype, MPI_Status *status)`
  - **Collectively** read/write using individual file pointer
- `MPI_File_sync(MPI_File fh)`
  - **Flush** all previous writes to the storage device

# Scientific Data Format: NetCDF & HDF5

- What is a Scientific Data Format?
  - A data format for **scientist** to store, access & operate their data **easily and efficiently**
- Key requirements:
  - **Self-Describing**: A file **includes information about the data** it contains.
  - **Portable**: A file can be accessed by computers with **different ways of storing integers, characters, and floating-point numbers**.
  - **Scalable**: Small subsets of large datasets in various formats may be **accessed efficiently through file interfaces, even from remote servers**.
  - **Appendable**: Data may be appended to a properly structured file **without copying the dataset or redefining its structure**.
  - **Sharable**: One writer and multiple readers may **simultaneously access the same file**.
  - **Archivable**: Access to all earlier forms of data will be **supported by current and future versions of the software**.

# Scientific Data Format: NetCDF & HDF5

- What is a Scientific Data Format?
  - A data format for **scientist** to store, access & operate their data **easily and efficiently**
- Key requirements:
  - **Sequential**: Data is stored in a sequential manner.
  - **Portable**: Data can be stored in a format that is independent of the hardware and software used to create it.
  - **Self-describing**: The data format is self-describing, meaning that the data can be accessed without copying the dataset or redefining its structure.
  - **Sharable**: One writer and multiple readers may **simultaneously** access the same file.
  - **Archivable**: Access to all earlier forms of data will be **supported by current and future versions of the software**.

Can you achieve these with a complex **doc file** or a **plain text file**?

# Scientific Data Format: NetCDF & HDF5

- Key features
  - A file contains its own directory (tree) structure
  - Each dataset is a multi-dimensional array
    - The dimension and size can be configured & changed
  - Each file entity (group & dataset) is self-describe
    - By its own metadata & attributes
  - The mapping between the dataset and disk layout can be controlled
    - Column or row major
- Visualization tools
  - HDFView
  - <https://www.hdfgroup.org/downloads/hdfview/>

