

Introduction to SIMD


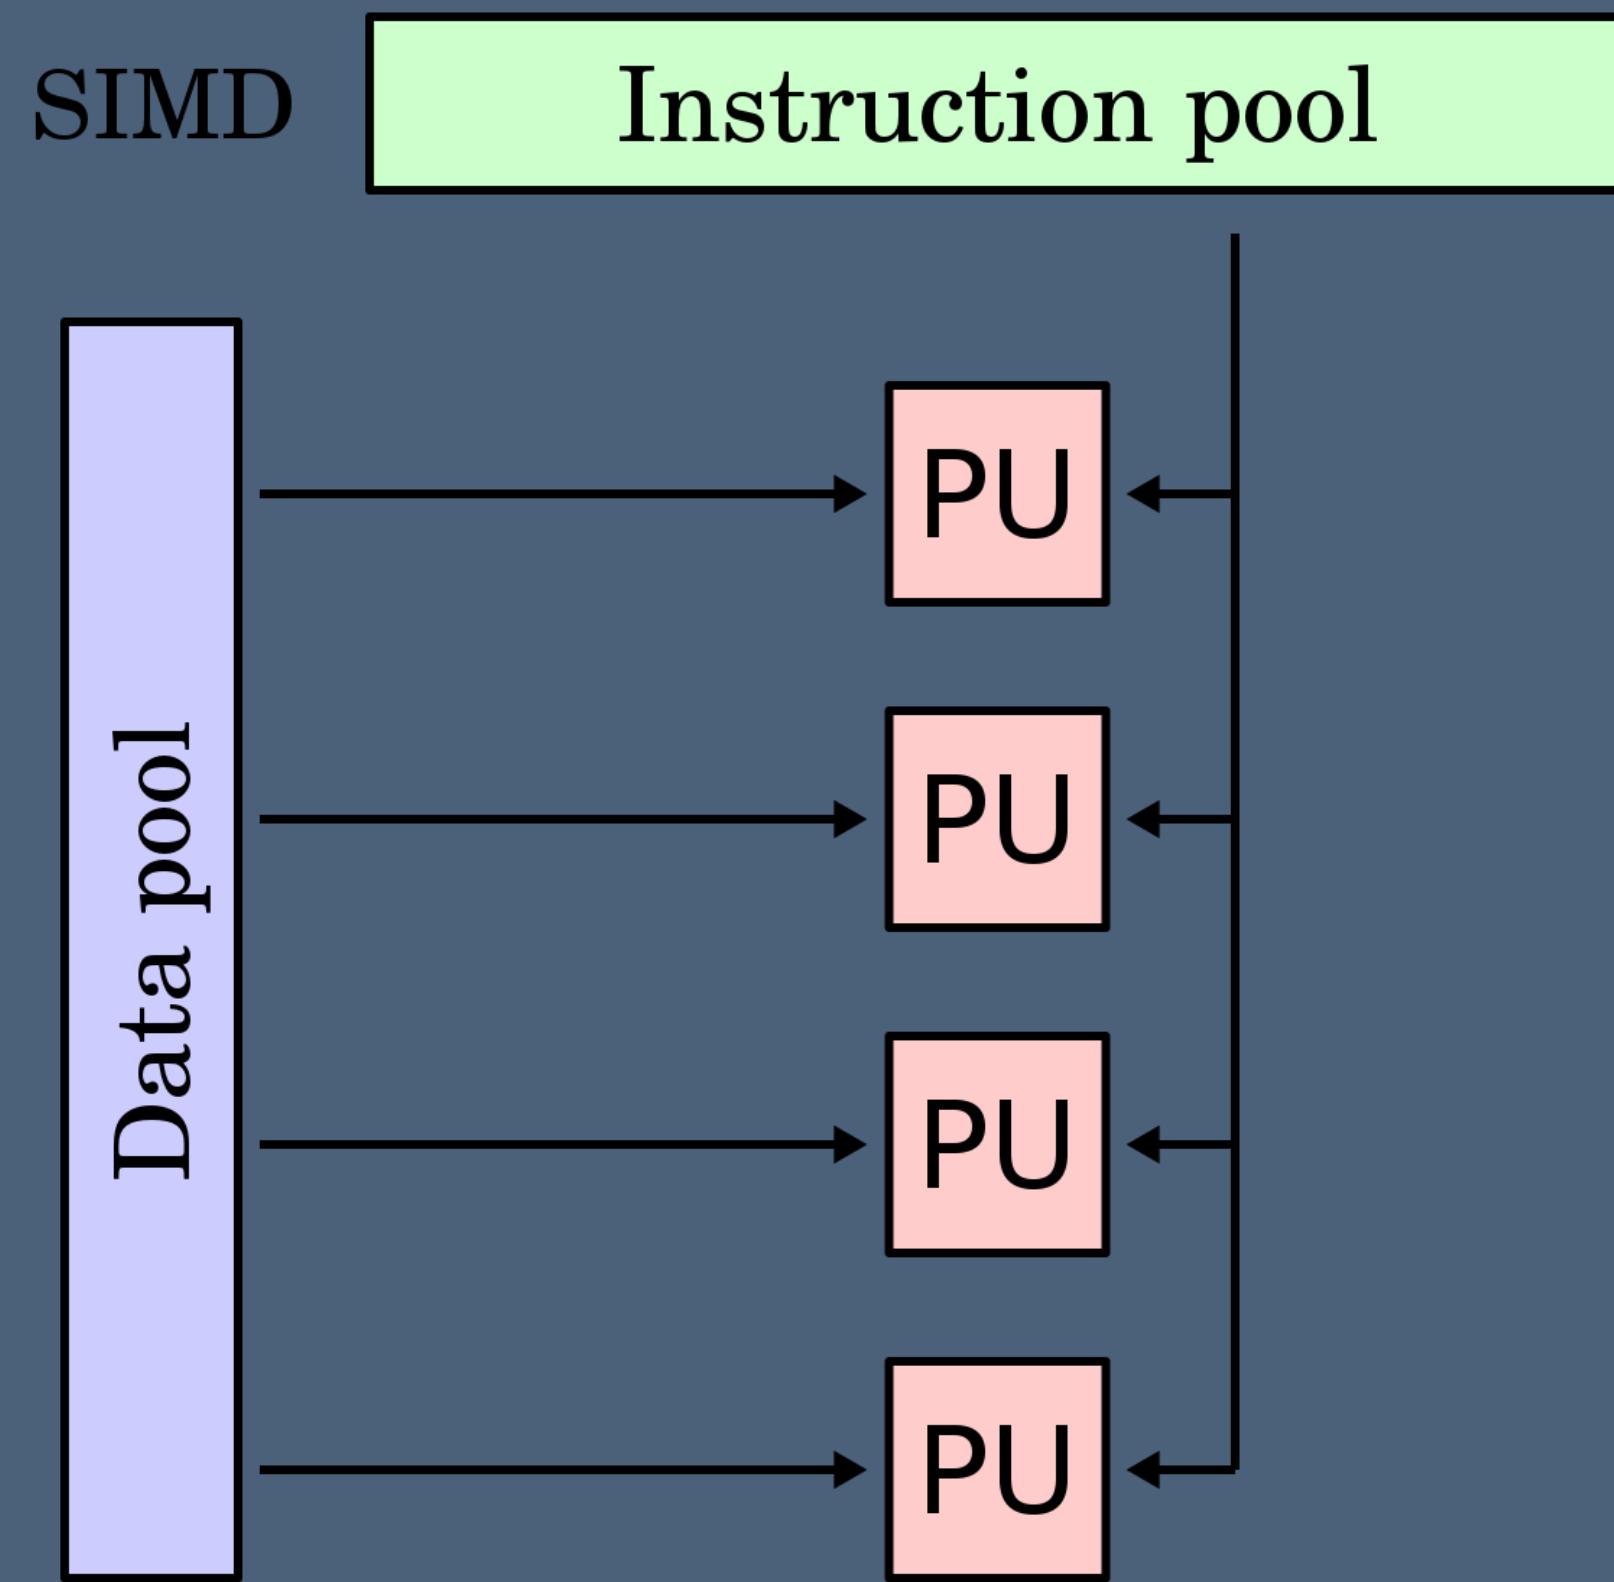
Fundamental High Performance Computing Cluster Practice

Background Ideas

- In most compute-intensive scenario, we often need apply the same operations to a series of data
- Nowadays, machine learning algorithms usually transform inputs to vectors
 - word2vec
 - wav2vec 2.0
 - Or even Life2Vec

SIMD

Single Instruction Multiple Data



Flynn's Taxonomy

Michael, 1966

	Single Data	Multiple Data
Single Instruction	SISD e.g., scalar CPU	SIMD e.g., vector processor, SIMD instructions, GPU
Multiple Instruction	MISD e.g., systolic array (Google's TPU)	MIMD e.g., multi-core CPU

Characteristics of Different SIMD Solutions

- Vector processor
 - Variable-length data
 - Ad-hoc application
 - Cray
- SIMD Instruction
 - Fix-length data
 - Generic application
- GPU
 - Multiple threads (SIMT, Single Instruction Multiple Thread)

Intel Intrinsic History

Intel SIMD ISA Evolution


SIMD extensions on top of x86/x87

64b
SIMD

128b
SIMD

256b
SIMD

512b
SIMD



MMX

Multi-Media Extension, 1996

- Pentium (32-bit)
- 64-bit width
- Integer-type operations only
- Reuse mantissa part of extended-precision floating-point in x87 FPU as registers

SSE & SSE 2, 3, 4.1, 4.2

Streaming SIMD Extension, 1999

- Pentium 3 (32-bit)
- 128-bit width
- Support single-precision floating-point (integer and double since SSE 2, 2001)
- Introduce 8 new dedicated XMM registers (16 registers in 64-bit processor)
- In fact, nowadays floating-point math adopts SSE

AVX & AVX 2

Advanced Vector Extension, 2011 & 2013

- Sandy Bridge (64-bit)
- 256-bit width
- Support integer, single- and double-precision floating-point
- Introduce 16 YMM registers, lower halves of which are treated as the original XMM ones

AVX 512

Advanced Vector Extension, 2016

- Skylake (64-bit)
- 512-bit width
- Support integer, single- and double-precision floating-point
- Introduce 32 ZMM registers, lower halves of which are treated as the original YMM ones

Criticisms & Issues

“Die a Painful Death?”

- Performance & clock rate, power consumption, die space
- Complex variants
- Not all instructions for 128-bit / 256-bit have corresponding ones for 256-bit / 512-bit

SIMD Programming

Fine-grained, Low-level

Functions that provide access
to other instructions without
writing assembly code.

What are Intrinsic?

Example

Population Count

- POPCNT, the CPUID flag / instruction came at the same time with SSE 4.2
- With GCC, one could use its `__builtin_popcount()` and `__builtin_popcountl()`
- The corresponding Intel Intrinsics are `_popcnt32()` and `_popcnt64()`

Check What CPU Flags are Supported

- Linux
 - `lscpu`
 - `cat /proc/cpuinfo`
- macOS
 - `sysctl -a machdep.cpu.features`

Compiler Flags

- GCC
 - `-m<target>`
- ICC
 - `-x<code>`

Header Files

```
#include <mmintrin.h> // MMX
#include <xmmintrin.h> // SSE
#include <emmintrin.h> // SSE2
#include <pmmintrin.h> // SSE3
#include <tmmintrin.h> // SSSE3
#include <smmintrin.h> // SSE4.1
#include <nmmintrin.h> // SSE4.2
#include <immintrin.h> // AVX, AVX2, FMA
```

Data Types

- 128-bit
 - `__m128`
 - `__m128d`
 - `__m128i`
- 256-bit
 - `__m256`
 - `__m256d`
 - `__m256i`
- 512-bit
 - `__m512`
 - `__m512d`
 - `__m512i`

Intrinsics Pattern

`_mm<width>_<operation>_<type>()`

- Width
 - 512, 256, or empty for 128
- Type
 - **ps** for packed single, **pd** for packed double; **epi8**, **epi32** and **epi64** for packed integer types operations, respectively
 - **si512**, **si256** and **si128** for operations on the whole integer register

Operation

Load / Store

- **load, store**
 - Load / store data between memory and register. Note that memory must be aligned (XMM requires 16-byte, YMM 32-byte, ZMM 64-byte)
- **loadu, storeu**
 - Load / store data between memory and register. Memory could be un-aligned.

Operation

Set

- **setzero**
 - Reset all bits in the register
- **set1**
 - Broadcast a value to all elements in the register
- **set**
 - Set values from lower to higher
- **setr**
 - Set values in reverse order

Operation

Extract

- **extract**
 - Get the element in the register

Operation

Arithmetic - Addition & Subtraction

- **add, sub**
 - Leave overflow as-is
- **adds, subs**
 - Integer addition and subtraction with “*saturation*”. Overflow results would be bound to maximum or minimum

Operation

Arithmetic - Multiplication & Division

- **mul**

- Note that for integers, only even-indexed elements would be extended and multiplied

- **mullo, mulhi**

- For integer multiplication, the results would be the lower / higher halves

- **div**

- Only for floating-point numbers

Operation

Arithmetic - Bitwise

- **and, or, xor**
- **slli, sll, sllv**
 - Shift left logically
- **srl, srl, srlv**
 - Shift right logically
- **srai, sra, sra**
 - Shift right arithmetically

Operation

Arithmetic - Comparison

- **cmpeq, cmpge, cmpgt, cmple, cmplt**
 - If the condition is true, all bits of the element in result would be set

Operation

Arithmetic - Math

- **min, max, avg, abs**
 - Average only available for unsigned integers (ceiling up)
- **sqrt, rcp, rsqrt**
 - Reciprocal ones are approximate

Operation

Shuffle & Permute

- **shuffle**
 - Select from 2 registers
- **permute**
 - Select from a register
- **blend, blendv**
 - Blend 2 registers

Intel Intrinsics Guide

<https://www.intel.com/content/www/us/en/docs/intrinsics-guide/index.html>

- Latency
 - The delay generated by the dependency-chain
- Throughput
 - Note that it is in CPI (cycle per instruction), which is reciprocal to common ways, IPC (instruction per cycle)

Case Study

<https://github.com/nevikw39/HPC-Vec>

Compute the Angle between 2 Vectors

- We have that $\mathbf{a} \cdot \mathbf{b} = |\mathbf{a}| |\mathbf{b}| \cos \theta$ and $|\mathbf{a}| = \sqrt{\mathbf{a} \cdot \mathbf{a}}$
- Thus we only need implement the dot product operation

Scalar Version

```
float dot(const float *const lhs, const float *const rhs)
{
    float y = 0;
    for (int i = 0; i < N; i++)
        y += i[lhs] * i[rhs];
    return y;
}
```

Vectorization with SSE

Naïve Approach

```
float dot_naive(const float *const lhs, const float *const rhs)
{
    assert(sizeof(float) == 4);
    float y = 0;
    for (int i = 0; i < N; i += 4)
    {
        __m128 li = _mm_load_ps(lhs + i), ri = _mm_load_ps(rhs + i);
        __m128 tmp = _mm_mul_ps(li, ri);
        _Alignas(16) static float arr[4];
        _mm_store_ps(arr, tmp);
        y += *arr + 1[arr] + 2[arr] + 3[arr];
    }
    return y;
}
```

Vectorization with SSE

Better Approach

```
float dot(const float *const lhs, const float *const rhs)
{
    assert(sizeof(float) == 4);
    __m128 y = _mm_setzero_ps();
    for (int i = 0; i < N; i += 4)
    {
        __m128 li = _mm_load_ps(lhs + i), ri = _mm_load_ps(rhs + i);
        __m128 tmp = _mm_mul_ps(li, ri);
        y = _mm_add_ps(y, tmp);
    }
    return y[0] + y[1] + y[2] + y[3];
}
```

Vectorization with AVX

```
float dot(const float *const lhs, const float *const rhs)
{
    assert(sizeof(float) == 4);
    __m256 y = _mm256_setzero_ps();
    for (int i = 0; i < N; i += 8)
    {
        __m256 li = _mm256_load_ps(lhs + i), ri = _mm256_load_ps(rhs +
i);
        __m256 tmp = _mm256_mul_ps(li, ri);
        y = _mm256_add_ps(y, tmp);
    }
    return y[0] + y[1] + y[2] + y[3] + y[4] + y[5] + y[6] + y[7];
}
```


Vectorization with AVX 512

```
float dot(const float *const lhs, const float *const rhs)
{
    assert(sizeof(float) == 4);
    __m512 y = _mm512_setzero_ps();
    for (int i = 0; i < N; i += 16)
    {
        __m512 li = _mm512_load_ps(lhs + i), ri = _mm512_load_ps(rhs + i);
        __m512 tmp = _mm512_mul_ps(li, ri);
        y = _mm512_add_ps(y, tmp);
    }
    return y[0] + y[1] + y[2] + y[3] + y[4] + y[5] + y[6] + y[7] + y[8] +
    y[9] + y[10] + y[11] + y[12] + y[13] + y[14] + y[15];
}
```

FMA

Fused Multiply-Add

- Perform rounding only once
 - Improve precision
- Leverage the pipeline
 - Improve performance
- Require AVX (even for 128-bit types)



Compute the Hamming Distance between 2 Strings

Count the number of indices that the values are not equal

- With vector instructions, we could compare several pairs of characters efficiently
- In light of the previous example, we should avoid reducing result too frequently
- For 8-bit integers, the result of equality comparison would be 0xFF (or -1) if equal or 0 otherwise
- If we accumulate the results, an 8-bit integer could maintain up to 256 times

Scalar Version

```
int hamming_dist()
{
    int y = N;
    for (int i = 0; i < N; i++)
        y -= i[str_a] == i[str_b];
    return y;
}
```

Vectorization with SSE

```
int hamming_dist()
{
    int y = N;
    __m128i x = _mm_setzero_si128();
    _Alignas(16) signed char arr[16];
    for (int i = 0; i < N; i += 16)
    {
        __m128i a = _mm_load_si128(str_a
+ i), b = _mm_load_si128(str_b + i);
        __m128i tmp = _mm_cmpeq_epi8(a,
b);
        x = _mm_add_epi8(x, tmp);
        if (!(i % (16 * 256)))
        {
            _mm_store_si128(arr, x);
            for (int j = 0; j < 16; j++)
                y -= (-j[arr] % 256 +
256) % 256;
        }
    }
    if (!(N % (16 * 256)))
    {
        _mm_storeu_si128(arr, x);
        for (int j = 0; j < 16; j++)
            y -= (-j[arr] % 256 + 256) %
256;
    }
    return y;
}
```


Vectorization with AVX 2

_mm256_cmpeq_epi8() requires AVX 2

```
int hamming_dist()
{
    int y = N;
    __m256i x = _mm256_setzero_si256();
    _Alignas(32) signed char arr[32];
    for (int i = 0; i < N; i += 32)
    {
        __m256i a =
        _mm256_load_si256(str_a + i), b =
        _mm256_load_si256(str_b + i);
        __m256i tmp =
        _mm256_cmpeq_epi8(a, b);
        x = _mm256_add_epi8(x, tmp);
        if (!(i % (32 * 256)))
        {
            _mm256_store_si256(arr, x);

```

```
        for (int j = 0; j < 32; j++)
            y -= (-j[arr] % 256 +
                256) % 256;
        x = _mm256_setzero_si256();
    }
}
if (!(N % (32 * 256)))
{
    _mm256_storeu_si256(arr, x);
    for (int j = 0; j < 32; j++)
        y -= (-j[arr] % 256 + 256) %
            256;
}
return y;
}
```



Find the Maximum Prefix Sum

- This example is far more complicated than previous ones as there are data dependencies
- Maximum could still be found by vectorized instructions
- For prefix sum, we could update by $2 \log_2 n$ vectorized instructions
- Note that the 256-bit shift and shuffle instructions are constrained by the 128-bit lane boundary

Scalar Version

```
int max_prefix()
{
    int y = 0;
    for (int i = 0, ps = 0; i < N; i++)
    {
        ps += i[arr];
        if (y < ps)
            y = ps;
    }
    return y;
}
```

Vectorization with SSE 4.1

_mm_max_epi32() requires SSE 4.1

```
int max_prefix()
{
    assert(sizeof(int) == 4);
    __m128i mx =
    _mm_setzero_si128(), sum =
    _mm_setzero_si128();
    for (int i = 0; i < N; i += 4)
    {
        __m128i tmp =
        _mm_load_si128(arr + i);
        tmp = _mm_add_epi32(tmp,
        _mm_slli_si128(tmp, 4));
        tmp = _mm_add_epi32(tmp,
        _mm_slli_si128(tmp, 8));
```

```
sum = _mm_add_epi32(sum,
tmp);


mx = _mm_max_epi32(mx, sum);
sum = _mm_shuffle_epi32(sum,
0b11111111);
    }
    mx = _mm_max_epi32(mx,
    _mm_slli_si128(mx, 4));
    mx = _mm_max_epi32(mx,
    _mm_slli_si128(mx, 8));
    return _mm_extract_epi32(mx, 3);
}
```

Vectorization with AVX 2

_mm256_max_epi32() requires AVX 2

```
int max_prefix()
{
    assert(sizeof(int) == 4);
    __m256i mx = _mm256_setzero_si256(), sum =
    _mm256_setzero_si256();
    for (int i = 0; i < N; i += 8)
    {
        __m256i tmp = _mm256_load_si256(arr + i);
        tmp = _mm256_add_epi32(tmp,
        _mm256_alignr_epi8(tmp,
        _mm256_permute2x128_si256(tmp, tmp,
        _MM_SHUFFLE(0, 0, 2, 0)), 16 - 4));
        tmp = _mm256_add_epi32(tmp,
        _mm256_alignr_epi8(tmp,
        _mm256_permute2x128_si256(tmp, tmp,
        _MM_SHUFFLE(0, 0, 2, 0)), 16 - 8));
        tmp = _mm256_add_epi32(tmp,
        _mm256_permute2x128_si256(tmp, tmp,
        _MM_SHUFFLE(0, 0, 2, 0)));
    }
```

```
sum = _mm256_add_epi32(sum, tmp);
mx = _mm256_max_epi32(mx, sum);
sum =
_mm256_set1_epi32(_mm256_extract_epi32(sum, 7));
}
mx = _mm256_max_epi32(mx,
_mm256_alignr_epi8(mx,
_mm256_permute2x128_si256(mx, mx, _MM_SHUFFLE(0,
0, 2, 0)), 16 - 4));
mx = _mm256_max_epi32(mx,
_mm256_alignr_epi8(mx,
_mm256_permute2x128_si256(mx, mx, _MM_SHUFFLE(0,
0, 2, 0)), 16 - 8));
mx = _mm256_max_epi32(mx,
_mm256_permute2x128_si256(mx, mx, _MM_SHUFFLE(0,
0, 2, 0)));
return _mm256_extract_epi32(mx, 7);
}
```

Auto-Vectorization

GCC


- **`-ftree-optimize`**
 - This is included since -O3 and above
- **`#pragma GCC ivdep`**
 - Tell the compiler that there is no dependency

OpenMP SIMD

Appears in OpenMP 4.0

- `#pragma omp simd`
- Compiler optimization should also be set

Wall Time



Packed Operation Percentage

